

SingGuard: Policy-Adaptive Multimodal Safeguarding with Dynamic Reasoning

SingGuard Team

¹AI Security Lab, Ant Group




Figure 1: Average F1 across six MLLM guard benchmark families spanning 35 underlying datasets.

Abstract

Vision-language models (VLMs) are now deployed across consumer, medical, financial, and enterprise applications, yet their expanding reach broadens the safety surface: risks may arise from text, images, or cross-modal composition, while moderation policies shift across products, regions, and deployment stages. Most existing guardrails either rely on fixed taxonomies or cover only a single modality, limiting adaptability when safety rules change at deployment time. We present **SingGuard**, a policy-adaptive multimodal guardrail model family that unifies text, image, image-text, and multilingual safety assessment within a single architecture. SingGuard treats the active policy as a runtime input: given natural-language rules, it performs rule-by-rule dynamic matching and predicts both the safety label and the specific triggered rule. To balance efficiency and interpretability, SingGuard supports multiple inference regimes along a fast-to-slow reasoning spectrum, from direct safety judgments to policy-grounded dynamic deliberation, and is further optimized with fast-slow decoupled reinforcement learning. We also introduce **SingGuard-Bench**, a multimodal guardrail benchmark with 58K+ examples spanning 80+ fine-grained risk types across text-only, image, image-text, attack, and dynamic-rule slices, including cross-modal joint-risk cases where each modality is individually harmless yet their composition implies unsafe intent. Across six benchmark families (35 datasets)—multimodal, image, text query, text response, multilingual query, and multilingual response safety—SingGuard achieves state-

of-the-art average F1 on every family. Dynamic-rule evaluation further shows improved policy-following accuracy from 0.6465 to 0.7415 under runtime policy shifts. Our code is available at <https://github.com/inclusionAI/SingGuard>.

1 Introduction

Vision-language models (VLMs) have become general-purpose multimodal foundation models (Hurst et al., 2024; Team et al., 2024; Google, 2024; Anthropic, 2024; Bai et al., 2025b; Chen et al., 2024; Bai et al., 2025a) and are increasingly deployed in consumer assistants, creative tools, and high-stakes domains such as medicine, finance, and enterprise decision support. This deployment trend substantially broadens the safety surface. Risk can appear in the user query or in the model response, arise from a single modality or from image-text composition, and depend on multilingual, multi-turn, or culturally specific context. At the same time, moderation policies are not fixed: they vary across products, regions, and deployment stages, and they change as new business rules or regulations are introduced. A practical guardrail must therefore go beyond single-modality classification under a static taxonomy. It should support heterogeneous multimodal inputs while conditioning its decisions on the active policy supplied at runtime.

Existing guardrails address only part of this requirement. Text guardrails such as Llama Guard, WildGuard, and Qwen3Guard achieve strong prompt or response moderation under predefined taxonomies (Inan et al., 2023; Han et al., 2024; Zhao et al., 2025; Zeng et al., 2024; Research, 2024; Kumar et al., 2025), but their policy space is typically a fixed label set with limited support for runtime extension. Multimodal guardrails such as Llama Guard 3 Vision, LlavaGuard, and GuardReasoner-VL extend moderation to images and image-text interactions (Chi et al., 2024; Zeng et al., 2025; Helff et al., 2025; Huang et al., 2025; Xu et al., 2025; Liu et al., 2025b). These systems demonstrate the importance of visual evidence, but most still assume static policy boundaries and classify risk directly rather than matching content against an open set of active rules. Recent policy-adaptive guardrails, including YuFeng-XGuard and policy-adaptive image guardrails, further show that fixed-policy training can overfit seen rules and degrade under unseen policy shifts (Lin et al., 2026; Piao et al., 2026). However, this line of work remains largely image-centered or text-centered and does not provide unified coverage of text, image, image-text, query-side, response-side, and multilingual moderation.

Another line of work introduces reasoning-based guardrails, such as GuardReasoner and GuardReasoner-VL, that generate rationales before making the final decision (Liu et al., 2025a;b). Reasoning is valuable not only for dynamic policies, but also for explanation, auditability, evidence grounding, and resolving borderline cases. However, a single always-on reasoning mode is not ideal for all moderation workloads. For most high-throughput moderation cases, the policy is stable and the input is unambiguous, so a direct first-pass judgment is already reliable enough while being much faster. Long chain-of-thought generation adds latency and often brings limited accuracy gain on these routine samples. Slow deliberation is most useful for cases where the expected label may change under a runtime policy, where the sample is counter-intuitive relative to the default taxonomy, or where the risk is implicit and must be inferred from cross-modal or multi-clause context. These heterogeneous needs motivate multiple inference regimes along a fast-to-slow reasoning spectrum: the guardrail should support direct judgments for clear cases, compact explanations when lightweight justification is sufficient, and policy-grounded rule-by-rule deliberation when the decision boundary is complex, changing, or audit-sensitive.

We introduce **SingGuard**, a policy-adaptive multimodal guardrail model family that unifies common moderation modalities and stages within a single framework. SingGuard uses a unified hierarchical safety taxonomy as the default rule space, while allowing the active policy to be extended at inference time. The policy can be provided as concise category names or as detailed natural-language rules specifying scope, exceptions, and domain-specific constraints. Given this policy, SingGuard performs rule-by-rule matching and predicts whether the content violates an active rule, rather than relying only on memorized taxonomy priors. This interface allows deployment teams to add or revise rules for domains such as health, finance, legal advice, privacy, or product-specific policies without retraining the model.

SingGuard supports multiple inference regimes along a fast-to-slow reasoning spectrum (Fig. 2). The *fast* mode directly emits a safety label and triggered category for low-latency moderation. The *dynamic* mode first makes an initial safe/unsafe judgment and then decides whether policy-grounded reasoning is needed. When stronger evidence or auditability is required, the *slow* mode performs explicit rule-by-rule deliberation over the active policy before producing the final answer. We train these behaviors with supervised fine-tuning over unified text and multimodal safety data, using a shared <fast> → <reasoning> → <answer> output grammar. We then apply fast-slow decoupled reinforcement learning: the first fast token is excluded from




Figure 2: Overview of **SingGuard**. Previous guardrails often cover only partial modalities, rely on fixed taxonomies or static policies, and use a fixed direct-answer or always-reasoning output path. SingGuard unifies multimodal inputs, open active policies, and dynamic reasoning paths in one policy-adaptive guardrail model.

the RL response update to reduce its anchoring effect on later reasoning, while the final answer is optimized with binary safety and fine-grained category rewards.

To provide a comprehensive multimodal guardrail evaluation, we introduce **SingGuard-Bench**, a multimodal guardrail benchmark with more than 80 fine-grained risk types organized under a three-level safety taxonomy. The benchmark covers text-only, image, and image-text samples; attack-style prompts such as typographic, semantic-isomorphic, narrative, and role-play attacks; and cross-modal joint-risk attacks where the image and text are individually benign but jointly imply unsafe intent. It also includes a dynamic-rule split that pairs each sample with both matching and non-matching runtime policies. SingGuard-Bench is constructed through a quality-oriented synthesis and filtering pipeline with three-level keyword expansion, multi-model consistency checking, and orthogonal dynamic-rule construction. We describe the construction process in Section 2.

As shown in Fig. 1, SingGuard achieves state-of-the-art average performance across six benchmark families: multimodal safety, image safety, text query safety, text response safety, multilingual query safety, and multilingual response safety. These families span 35 underlying datasets and evaluation splits. The results indicate that a single policy-adaptive guardrail can deliver broad performance across modalities, languages, and moderation stages, rather than improving only on an isolated benchmark. In dynamic-policy evaluation, SingGuard-slow improves average accuracy from 0.6465 (Qwen3-VL-8B) to 0.7415, showing that runtime rule conditioning helps the model follow changing active policies instead of relying solely on training-time taxonomy priors. Full per-dataset results are reported in Section 4.

Our contributions are summarized as follows:

- We propose SingGuard, a policy-adaptive vision-language guardrail model family that unifies text, image, image-text, multilingual, query-level, and response-level moderation under an open runtime-policy interface.
- We design a flexible fast-to-slow inference scheme that supports direct judgment, adaptive dynamic reasoning, and evidence-oriented slow deliberation, together with a fast-slow decoupled reinforcement learning objective that preserves low-latency fast judgments while reducing their anchoring effect on policy-grounded reasoning.
- We build SingGuard-Bench, a comprehensive MLLM guardrail benchmark covering harmful recall, benign-sensitive precision, attack robustness, modality composition, keyword coverage, reasoning depth, multilingual safety, and dynamic rules.
- We show that SingGuard reaches state-of-the-art average performance across six benchmark families spanning 35 underlying datasets, while also improving dynamic-policy accuracy under changing active rules.

2 Method

Method overview. SingGuard is a policy-conditioned multimodal guardrail model: given target content and an active policy, it predicts the safety label, the triggered active rule, and, when needed, a policy-grounded reasoning trace. The method addresses three limitations of static guardrails. First, fixed taxonomies fail when deployment policies change, so we define a runtime-policy interface over a unified hierarchical safety taxonomy and construct policy-conditioned examples where the same content can receive different labels under different active rules. Second, ordinary label supervision can entangle the short initial judgment with later deliberation, so we combine cold-start SFT with fast-slow decoupled DAPO to reduce the anchoring effect of the first fast token on subsequent reasoning. Third, a fixed reasoning depth is inefficient: fast mode returns a direct decision, dynamic mode first makes an initial safe/unsafe judgment before deciding whether further policy-grounded reasoning is needed, and slow mode performs explicit rule-by-rule deliberation for evidence- or audit-sensitive cases. The rest of this section presents the task formulation, policy taxonomy, data construction, SFT stage, RL stage, dynamic inference modes, and model distillation.

2.1 Task Formulation

We formulate policy-adaptive safety classification as an instruction-following task. The model is given explicit moderation instructions, an active safety policy, target content, and output-format requirements, and is expected to generate a structured moderation decision. Unlike a fixed-label classifier, the decision must be grounded in the policy supplied at inference time.

Formally, an input instance is represented as

$$x = (q, I, a), \quad (1)$$

where q denotes the user’s query, I denotes zero or more images, and a denotes the assistant response when response-level moderation is required. This formulation covers text-only queries, image inputs, image-text queries, assistant responses, and full query-response conversations. Depending on the evaluation setting, SingGuard can assess the user query, the assistant response, or the query-response pair under the same policy-conditioned interface.

Given an active policy $P = \{r_1, \dots, r_n\}$, SingGuard solves

$$f_\theta(x, P) \rightarrow (y, z, c), \quad (2)$$

where $y \in \{\text{safe}, \text{unsafe}\}$ is the overall assessment, z is an optional policy-grounded reasoning trace, and $c \in \mathcal{T}(P) \cup \{\text{Safe}\}$ is the final triggered category or rule title. Here $\mathcal{T}(P)$ denotes the set of valid category names or rule titles explicitly activated by the current policy.

For query assessment, the model evaluates whether the user’s latest query, together with any associated image and dialogue context, violates the active policy. For response and query-response assessment, the model judges whether the assistant response provides harmful assistance, unsafe confirmation, privacy leakage, or other policy-violating content in context, while jointly considering user intent, visual evidence, assistant behavior, and cross-modal interaction.

The output schema is introduced in Section 2.4, with complete prompt templates provided in Appendix A. In all modes, the final answer is bound to the active policy P : uncovered risks may still be labeled safe, while newly introduced active rules must be enforced.

2.2 Safety Policy and Taxonomy

A robust safety policy is the conceptual core of a production-grade guardrail. To be effective, the policy must be comprehensive enough to cover diverse harm types, systematic enough to provide a consistent classification framework, and fine-grained enough to support precise risk attribution rather than undifferentiated binary labels.

SingGuard therefore uses a unified hierarchical safety taxonomy as its built-in default safety policy, as summarized in Table 14. The taxonomy contains 8 primary dimensions, 27 secondary categories, and more than 80 fine-grained risk types. It covers sexual content, real-world crime and public safety, unethical behavior, cybersecurity and information manipulation, agent safety, politically sensitive content, animal abuse, and benign content. This hierarchical structure provides a stable policy space for aligning heterogeneous text, image, image-text, and query-response safety data under a shared annotation schema.

The categories in Table 14 define the out-of-the-box policy that SingGuard can apply directly. However, the policy boundary is not fixed to this taxonomy. At inference time, the active policy can be instantiated as the




Figure 3: Policy-conditioned data and training pipeline. SingGuard aligns heterogeneous safety data with active policies, augments them with policy views and counterfactual labels, verifies semantic consistency, and uses the resulting supervision for fast and slow guardrail training.

full taxonomy, a task-specific subset, a narrowed or expanded version of existing rules, or newly introduced domain-specific rules. This design allows operators to adapt the guardrail to product requirements, regional constraints, and emerging risks while still requiring the model to ground each decision in the currently active rule set.

2.3 Data Collection and Annotation

Policy-adaptive moderation requires data that are diverse in modality, stable in safety semantics, and explicitly coupled with the policy shown at inference time. We organize the annotation pipeline around four complementary corpora, all aligned to the hierarchical safety taxonomy from Section 2.2: (i) re-annotated open-source safety data for broad risk coverage, (ii) policy-grounded synthetic data that fill long-tail and cross-modal gaps, (iii) dynamic rule-conditioned data that bind the decision to the active policy rather than a memorized taxonomy, and (iv) chain-of-thought (CoT) supervision that teaches rule-by-rule verification. The final corpus contains roughly 2.5M textual samples and 600K multimodal samples, of which more than 1M carry CoT traces. All samples follow a unified format containing the user query, optional image(s), optional assistant response, and dialogue context. Query-side and response-side labels are kept separately, allowing the same data source to support prompt screening, response checking, and joint query–response moderation.

2.3.1 Integration and Re-annotation of Open-Source Data

We use public safety benchmarks and training sets as the foundation of our corpus. Since different sources follow different annotation protocols and risk granularities, directly merging them would introduce inconsistent supervision. We therefore design an LLM-driven re-annotation pipeline that normalizes all data into our unified risk taxonomy and assigns each accepted instance both a risk label and a short rationale as auxiliary supervision. The pipeline covers harmful prompts, benign but sensitive prompts, harmful responses, safe refusals, and helpful safe completions on the text side, and incorporates AIGC, internal multimodal data, BeaverTails-V, MMDS, UnsaferBench, and related multimodal sources covering visual jailbreaks, image–text intent composition, multimodal query–response pairs, and bottom-line visual risks such as NSFW content, hateful memes, weapons, crime scenes, and violence.

For open-source data with existing annotations, we first map the original label space to our risk taxonomy using predefined category correspondences. Each mapped sample is then passed to an LLM or MLLM judge

for independent verification under our risk definitions. If the model confirms the mapping, the sample is retained; otherwise, it is re-annotated according to the safety annotation guideline. Ambiguous cases, including samples from policy gray zones or those with possible safe/unsafe polarity flips, are not directly merged into the corpus but are instead treated as hard cases and moved into the unlabeled-data annotation pool.

For unlabeled data and hard cases, we apply a customized safety annotation guideline and use multiple open-source LLMs/MLLMs, such as Qwen3.5-397B-A17B, KIMI-K2.6, and GLM4.5V, for automatic annotation. We combine self-consistency within each model with agreement across models, and retain a sample only when the models produce consistent judgments. All accepted samples further pass a two-level consistency check: L0 verifies the binary safe/unsafe decision, and L1 verifies the fine-grained risk category. Samples failing L0 are discarded, while those passing L0 but failing L1 are routed back as hard cases or synthesis-side seeds for dynamic data construction (Section 2.3.3). On an internal human-annotated validation set, this ensemble annotation pipeline achieves over 0.9 accuracy on safety-level labels.

2.3.2 Policy-Grounded Data Synthesis

Open-source data are long-tailed across policy categories and weakly cover cross-modal cases, so we build a policy-grounded synthetic pipeline that decomposes the safety policy into fine-grained categories and, per category, expands a small set of high-quality seeds through prompt, response, multi-category, multimodal, attack, and multilingual synthesis, with a unified secondary verification at the end.

Text-side prompt and response synthesis. For each fine-grained risk category, we extract policy keywords and use them to synthesize aligned text-side data. A red-teamed jailbreak model generates harmful prompts and responses, while an aligned model generates harmless contrastive prompts, policy-compliant refusals, and helpful safe completions. We also sample multiple policy definitions to create multi-category composites, and retain only samples whose independent re-annotation matches the intended labels.

Multimodal and cross-modal attack synthesis. For multimodal data, we use policy keywords to generate or retrieve harmful and benign images and texts, then combine them into image-text pairs. The accepted pairs cover different modality polarities, including image-unsafe/text-safe, image-unsafe/text-unsafe, image-safe/text-unsafe, image-safe/text-safe, and image samples. We further build cross-modal attacks by splitting a policy-relevant harmful intent across image and text, so that each modality appears benign alone but their composition becomes unsafe.

Multilingual extension and secondary verification. We extend the Chinese and English text-side data to additional languages while preserving the original policy schema. Because translation may weaken harmful intent or shift the safety label, each translated sample is re-verified under the same policy definition. If the inferred intent is inconsistent with the target risk category, the sample is filtered out, ensuring that translated prompts, responses, and multi-category composites remain semantically faithful and policy-consistent.

2.3.3 Dynamic Rule Data Construction

Static guardrails often learn a shortcut between surface content patterns and fixed safety labels. As a result, simply placing a new policy block in the prompt does not guarantee that the model will adapt its judgment to the new policy. To address this limitation, SingGuard constructs dynamic rule-conditioned examples in which the same content is evaluated under different active policies. The correct label is always recomputed according to the currently provided rules, so the model must ground its decision in policy semantics rather than memorized category names.

Policy-view diversification. For each filtered sample, we instantiate multiple views of the active policy. These include **full rules**, where the complete safety taxonomy is provided; **subset rules**, where only part of the taxonomy is activated; **single-rule views**, where only one relevant or irrelevant rule is shown; and **merged or rewritten rules**, where several rules are compressed into one compact description or rephrased with different wording and titles. We also vary the presentation style of the rules, including full-text rules, summaries, title-only items, plain numbering, and dynamically lettered titles. These views prevent the model from relying on fixed rule order, prompt length, or category titles, and require the final answer to match the currently active policy.

Counterfactual supervision and policy-shift augmentation. The supervision label is recomputed against the active policy rather than copied from the original safety taxonomy. To create policy-shifted supervision, we edit existing base rules by expanding the rule scope, adding exemption clauses, broadening or narrowing decision boundaries, changing evidence or intent requirements, and adjusting rule priority. The same anchor

sample is then relabeled under the edited policy, yielding **unsafe**→**unsafe**, **unsafe**→**safe**, **safe**→**unsafe**, and **safe**→**safe** transitions. For example, an unsafe sample can become safe when the matched rule is narrowed or an exemption is added, while a previously safe sample can become unsafe when the active policy expands the covered boundary. These counterfactual labels force the model to follow the current rule text rather than a fixed taxonomy prior.

Rule generation and verification. Beyond edits to existing categories, we also generate new dynamic audit rules outside the original taxonomy according to the sample content. These rules are written as general policy constraints rather than sample-specific descriptions, such as prohibiting certain behavior in a sports scenario, restricting unsafe claims in a legal setting, or disallowing particular advice in health, finance, or privacy contexts. A teacher LLM proposes the new rule and its intended safe/unsafe outcome for the anchor sample, and verifier models check that the rule is fluent, non-duplicative, sufficiently general, and semantically applicable. Before training, every dynamic example undergoes a final consistency check: the active rule must match the sample when the label is unsafe, the recomputed label must agree with the edited or newly added policy, and the final `<answer>` must be parseable and match an active rule title. Inconsistent samples are discarded or returned for rule revision, while accepted samples are used with the unified training schema described in Section 2.4.

2.3.4 Chain-of-Thought Reasoning Data Construction

To equip SingGuard with the reasoning and rule-by-rule verification ability required by the hybrid reasoning interface, we apply a large-scale CoT annotation pass on top of the fast-judgment data from Sections 2.3.1–2.3.3, producing more than 1M reasoning examples. Annotation is performed in a *rule-grounded, one-rule-at-a-time* manner under the full safety taxonomy together with its dynamic-rule extensions (matched and unmatched rules inherited from Section 2.3.3): the teacher walks through the active policy scope rule by rule, judges each rule as hit / not-hit / not-applicable with explicit evidence, and aggregates the per-rule judgments into the final `<answer>`—the first hit rule yields *unsafe* with its rule title, while *safe* is emitted only after the entire policy scope has been exhausted with no rule firing.

Hierarchical rule-grounded annotation. For the base safety categories, CoT annotation uses a richer teacher prompt than the prompt used during ordinary training. Given a case, the teacher receives the full rule set together with detailed rule descriptions, sub-clauses, examples, and exceptions, and is asked to judge the case rule by rule. This produces traces that expose the internal knowledge of each rule—from top-level category, to sub-clause, to supporting evidence—rather than only a final category label. After the traces are accepted, they can be paired during training with simpler and more diverse active-policy prompts, such as summarized rules, subset rules, merged rules, or title-focused prompts. This lets the model learn detailed rule-grounded reasoning while remaining robust to different policy presentations at inference time.

Lightweight cleansing and two-stage consistency acceptance. After annotation finishes, raw outputs first go through a lightweight deterministic pass that checks structural validity (tag well-formedness, length sanity) and applies a wording-normalization step on the reasoning core. Cleaned traces then go through a two-stage consistency acceptance pipeline. The first stage is a *CoT consistency* check: the rule-by-rule reasoning, the opening safe/unsafe judgment, and the supporting evidence must be mutually coherent; samples that fail this check are re-routed for re-annotation rather than discarded, so high-value but format-flawed cases are recovered. The second stage is an *answer consistency* check: the final `<answer>` category or rule title must exactly match the ground-truth hierarchical label; samples that still disagree after re-annotation are dropped. The accepted CoT examples are then merged with the fast-judgment data and used jointly during cold-start SFT (Section 2.4).

2.4 Stage 1: Policy-Conditioned Cold-Start SFT

The first stage is a policy-conditioned cold start whose goal is to inject the *judge–reason–review* paradigm into the base VLM under a shared output grammar, rather than to make every category appear with a long reasoning trace.

Unified training schema. All SFT samples are normalized under the same active-policy interface: the input contains the active policy and target content, and the output follows one of two compact forms. The *fast* form begins with a plain safe or unsafe token and then gives the final active-rule title in `<answer>`. The *fast-slow* form keeps the same opening decision and adds a `<reasoning>` span before the final answer, where the model summarizes the content, checks the active rules, and then produces the reviewed judgment. When multiple active rules are triggered or response-level refusal judgment is required, each item is written on a separate line inside `<answer>...</answer>`. Full prompt templates are provided in Appendix A.

Rule-by-rule policy grounding. Because the active policy can change at inference time, the reasoning span is not a free-form rationale but a structured rule-by-rule policy match. Given an active policy $P = \{r_1, \dots, r_n\}$, the reasoning is decomposed as $z = (z_{\text{summary}}, m_1, \dots, m_n, z_{\text{final}})$:

- z_{summary} describes the content (text, image, image–text relation, and response context);
- each $m_i \in \{\text{hit}, \text{not-hit}, \text{not-applicable}\}$ is the per-rule judgment for r_i , grounded in observable evidence;
- z_{final} aggregates the rule checks into the final decision.

This decomposition makes the trace auditable: an *unsafe* verdict must expose the triggered rule and supporting evidence, and a *safe* verdict must show that no active rule fires rather than silently relying on an inactive taxonomy prior.

Data mixture and training objective. Stage 1 mixes a large pool of fast-judgment samples with a smaller set of verified CoT-reasoning samples from Section 2.3.4. The output mode is controlled by the training prompt: some samples supervise a direct fast answer, while CoT-annotated samples supervise the fast-slow format with an explicit reasoning span. The fast samples provide broad coverage over modalities, policies, and risk categories, while the CoT subset teaches a reusable rule-by-rule reasoning procedure. Because the reasoning trace is grounded in the active policy rather than tied to a fixed sample template, this ability can generalize to cases where no explicit CoT annotation is available.

We optimize a single autoregressive log-likelihood with three field-level terms. Let X denote the full input (active policy plus target content), y_{cls} the opening safe/unsafe token, y_{reason} the <reasoning> span, and y_{cat} the final <answer> category. The objective is:

$$\mathcal{L}_{\text{SFT}} = \lambda_{\text{cls}} \mathcal{L}_{\text{cls}}(y_{\text{cls}} \mid X) + \lambda_{\text{reason}} \mathcal{L}_{\text{reason}}(y_{\text{reason}} \mid X, y_{\text{cls}}) + \lambda_{\text{cat}} \mathcal{L}_{\text{cat}}(y_{\text{cat}} \mid X, y_{\text{cls}}, y_{\text{reason}}), \quad (3)$$

where $\mathcal{L}_{\text{reason}}$ is masked out for fast-only samples and $\lambda_{\text{cls}}, \lambda_{\text{reason}}, \lambda_{\text{cat}}$ are field-level loss weights. This field-level supervision trains direct judgment and policy-grounded reasoning under the same output grammar, and Stage 2 further refines when and how the reasoning path should revise the initial decision.

2.5 Stage 2: Fast–Slow Decoupled DAPO

Most moderation cases can be handled by a direct judgment, but a small portion of hard examples require explicit reasoning. These include newly added dynamic rules, matched or unmatched distractor rules, policy-shifted counterfactuals, boundary cases near an exemption, and multimodal samples whose risk depends on the interaction between image and text. After cold-start SFT, we use the SFT checkpoint as the initial policy and optimize SingGuard with DAPO on a reasoning-heavy data mixture biased toward these dynamic-rule cases rather than replaying the full SFT mixture uniformly. SFT teaches the model the shared <fast> $\rightarrow$ <reasoning> $\rightarrow$ <answer> grammar, but it can also make the first fast judgment overly sticky: once the model emits an initial safe/unsafe token, the later reasoning span may simply rationalize that token instead of independently checking the policy and correcting the final answer. Stage 2 therefore preserves the fast path for easy cases while weakening its anchoring effect on the slow reasoning path, especially on policy-shifted examples where the default taxonomy prior is unreliable.

For each training query and active policy, the current model samples a group of G candidate outputs:

$$\mathcal{O} = \{o_1, \dots, o_G\}, \quad o_i = (y_{\text{cls}}^{(i)}, z^{(i)}, y_{\text{ans}}^{(i)}). \quad (4)$$

Each output follows the same <fast> $\rightarrow$ <reasoning> $\rightarrow$ <answer> format used in SFT. The group contains diverse initial judgments, reasoning traces, and final answers. We score each completed answer with a two-level safety reward:

$$R = 0.8 R_{\text{binary}} + 0.2 R_{\text{category}}, \quad (5)$$

where R_{binary} rewards safe/unsafe correctness and R_{category} rewards exact active-rule correctness in the final answer field. The binary term is dominant because moderation first requires the correct safety polarity: a safe sample should remain safe, while any unsafe category should be recognized as unsafe. The category term then encourages the reviewed answer to select the matched active rule or Safe. We additionally track whether the answer field is parseable, but the training reward in this stage is determined by the binary and category components.

To decouple fast and slow thinking during policy optimization, we mask the first generated response token after reward evaluation. Concretely, rollout first produces the full response under the SFT schema, the reward function decodes the complete valid response and places the scalar outcome reward on the last valid response token. The trainer then applies a prefix mask to the response mask, setting the first response position to zero

before DAPO advantage computation and actor update. Therefore the group-normalized outcome reward is still computed from the complete fast–slow answer, but the first generated token receives neither advantage nor policy-gradient loss. In DAPO, the normalized sequence score is broadcast only to unmasked response positions; in the actor objective, the clipped policy loss and optional KL loss are averaged with the same response mask. This makes the initial fast token act as a low-latency prefix rather than the sole optimization anchor for the whole response. The subsequent <reasoning> and <answer> tokens are still optimized by the group-relative reward, so a sampled trajectory can receive credit when the slow path corrects an unreliable initial judgment and produces the right final category.

Following DAPO, rewards are normalized within the sampled group to obtain an advantage for each output:

$$A_i = \frac{R_i - \text{mean}(\{R_j\}_{j=1}^G)}{\text{std}(\{R_j\}_{j=1}^G) + \epsilon}. \quad (6)$$

The model is then updated to increase the likelihood of outputs with higher advantages and suppress outputs with lower advantages. This makes optimization depend on relative quality among multiple sampled fast–slow trajectories rather than on a single deterministic completion. In practice, the update encourages trajectories whose final <answer> preserves the correct safety polarity and category even when the initial fast judgment is uncertain.

For hybrid reasoning, this stage is particularly important. The initial decision is useful for low-latency moderation, but the subsequent reasoning must still be able to verify or revise it when the active policy is ambiguous, newly introduced, or conflicts with the model’s memorized taxonomy prior. By masking the first fast token from the RL update while rewarding the final <answer>, Stage 2 turns the fast field into a provisional decision and gives the slow path room to recover from fast-thinking mistakes.

2.6 Policy-Grounded Dynamic Inference

After SFT and fast–slow decoupled DAPO, SingGuard uses the same trained fields under three inference modes. This separates the training format from the deployment budget: the model learns both compact decisions and policy-grounded reasoning during training, while deployment can choose how much of the sequence to decode.

Fast judgment. The fast mode directly emits the opening safe/unsafe decision and the <answer> field from the unified schema in Section 2.4. It is used for low-latency online moderation and high-throughput benchmark evaluation. For deployments with many active rules, we further support RI-Mask parallel rule checking (Appendix C), which preserves one-rule-at-a-time isolation while sharing the common content prefix. Since the output is short and parseable, it also provides the first decision used by the adaptive router.

Policy-grounded reasoning. The slow mode treats the initial decision as provisional, analyzes the content against the active policy in the <reasoning> field, and then emits a reviewed <answer>. This mode is used for ambiguous, audit-sensitive, or dynamically changing policies.

Hybrid reasoning. The hybrid mode implements adaptive early exit. The model first decodes only the initial binary safety decision, safe or unsafe. If this label is valid and sufficiently confident, decoding terminates immediately. Otherwise, the model continues with policy-grounded verification and emits a reviewed answer. Thus, easy cases behave like direct classification, while difficult cases receive the benefit of slow reasoning without requiring a separate mode-switching model.

Formally, let $y_0 \in \{\text{safe}, \text{unsafe}\}$ denote the initial binary decision. We compute confidence by normalizing only over the two safety-label probabilities:

$$s(y_0) = \frac{p_\theta(y_0 \mid x, P)}{p_\theta(\text{safe} \mid x, P) + p_\theta(\text{unsafe} \mid x, P)}, \quad (7)$$

where x is the target content and P is the active policy. The initial decision is accepted only when

$$y_0 \in \{\text{safe}, \text{unsafe}\} \wedge s(y_0) \geq \tau, \quad (8)$$

where τ controls the accuracy–latency trade-off. If the normalized confidence is below τ , SingGuard continues generation to produce the policy-grounded reasoning trace and reviewed answer.

2.7 Model Distillation via On-Policy GKD

After Stage 2, the 8B SingGuard model provides stronger policy-grounded reasoning and more stable dynamic-rule judgments. In contrast, the 2B model shows a larger gap on the cases that depend most on

reasoning capacity, including newly introduced dynamic rules, boundary rules with exemptions, long-tail risk categories, and samples where the reasoning trace must revise an initial fast judgment. The remaining gap is therefore not only a label-coverage problem, but a behavior-alignment problem: the compact model must learn how the stronger model interprets active rules and recovers from its own uncertain generations.

We address this gap with on-policy distillation (OPD) under the generalized knowledge distillation framework (Agarwal et al., 2024). For each prompt x drawn from the policy-conditioned mixture, the 2B student first samples its own response $\hat{y}$ under the same fast-slow output schema. The frozen 8B teacher then provides token-level target distributions on the student’s generated prefix. The student is optimized to match the teacher distribution with a bidirectional distillation objective:

$$\mathcal{L}_{\text{GKD}} = \mathbb{E}_{x \sim \mathcal{D}, \hat{y} \sim \pi_{\theta}(\cdot | x)} \left[\alpha D_{\text{KL}} \left(\pi_{\phi}^{\text{teacher}}(\cdot | x, \hat{y}_{<t}) \parallel \pi_{\theta}(\cdot | x, \hat{y}_{<t}) \right) + (1 - \alpha) D_{\text{KL}} \left(\pi_{\theta}(\cdot | x, \hat{y}_{<t}) \parallel \pi_{\phi}^{\text{teacher}}(\cdot | x, \hat{y}_{<t}) \right) \right]. \quad (9)$$

where π_{θ} is the 2B student, $\pi_{\phi}^{\text{teacher}}$ is the frozen 8B teacher, and α balances the teacher-to-student KL and the student-to-teacher reverse KL. The forward direction encourages the student to cover the teacher’s high-probability decisions, while the reverse direction gives a sharper on-policy signal on tokens sampled from the student’s own state space. Since the distillation target is computed on student-written responses rather than teacher-written traces, training directly exposes the teacher to the student’s typical errors on dynamic rules, edge cases, and long-tail categories. This aligns the training distribution with the student’s inference distribution and helps the compact model learn how to repair weak reasoning paths instead of merely imitating polished teacher outputs.

This OPD stage improves the 2B model’s dynamic-policy following and policy-grounded reasoning ability while preserving the low-latency deployment footprint; the quantitative effect is reported in Section 4.

3 SingGuard-Bench

Existing multimodal safety benchmarks provide valuable coverage of visual jailbreaks, image safety, preference alignment, or dialogue safety, but they leave several gaps for production moderation. First, their category systems often stop at broad labels, with limited subcategory definitions, sparse keyword coverage, and incomplete per-category coverage, making fine-grained error analysis unreliable. Second, many benchmarks contain limited benign-sensitive content, so a model can appear strong by over-blocking rather than by separating harmful content from legitimate medical, artistic, educational, or security-research material. Third, most datasets use static labels under a fixed policy and therefore cannot test whether a guardrail follows a newly supplied rule. Finally, common attack suites rarely include strict cross-modal hidden intent, where the image is safe alone and the text is safe alone, but their composition expresses an unsafe request.

We introduce **SingGuard-Bench**, a policy-conditioned multimodal safety benchmark for evaluating guardrail models. The benchmark contains 56,340 test examples, including 40,663 image samples, 13,677 multimodal samples, and 2,000 dynamic-rule examples split across image and text tracks. It spans more than 80 fine-grained risk types under the unified taxonomy in Section 2.2, and covers bottom-line visual harms, ordinary harmful image-text cases, benign-sensitive negatives, adversarial transformations, and cross-modal hidden-intent examples.

Each instance is normalized into a common guardrail interface with target content, an active policy, a safety label, and a fine-grained risk category or rule title. For dynamic-rule examples, the same content can be paired with matching or non-matching policy rules, so the correct label is defined by the active rule set rather than by a static dataset label. This design supports a unified evaluation of harmful-content recall, precision on benign-sensitive cases, cross-modal intent reasoning, and policy-following behavior. Table 1 situates SingGuard-Bench among representative multimodal safety benchmarks.

Evaluation axes. SingGuard-Bench is organized around four complementary axes:

- **Sample type:** adversarial attack samples, direct harmful samples, benign-sensitive samples, and dynamic-rule samples, corresponding to adversarial robustness, ordinary harmful recall, precision control, and policy adaptation.
- **Risk-category and keyword coverage:** the A–G unsafe taxonomy together with keyword-level semantic coverage for each fine-grained rule, reducing hidden category holes across long-tail safety risks.

Table 1: SingGuard-Bench comparison.

Benchmark	Test set scale	Image-only data	Image-text risk	Cross-modal hidden intent	Dynamic policy rules	Fine-grained taxonomy
MM-SafetyBench (Liu et al., 2023)	5,040	✗	✓	✗	✗	Limited
UnsafeBench (Qu et al., 2024)	2,040	✓	✗	✗	✗	✓
VLGuard (Zong et al., 2024)	1,000	✗	✓	✗	✗	Limited
JailBreakV-28K (Luo et al., 2024)	28,000	✗	✓	✗	✗	Limited
SPA-VL (Zhang et al., 2024)	7,530	✗	✓	✗	✗	✓
BeaverTails-V (Ji et al., 2025)	1,180	✗	✓	✗	✗	✓
MMDS (Huang et al., 2025)	4,484	✗	✓	Partial	✗	✓
VLSBench (Hu et al., 2024)	2,241	✗	✓	✗	✗	Limited
SingGuard-Bench (ours)	56,340	✓	✓	✓	✓	80+ risk types

- **Attack method:** cross-modal hidden-intent data, where a harmful intent is split across the image and text modalities so that neither modality is sufficient alone, together with conventional transformations such as Typography and Patch Shuffle.
- **Dynamic policy coverage:** per-sample policy sets containing matching and non-matching rules, designed to test whether the model follows the active policy instead of memorizing a fixed taxonomy.

3.1 Construction Pipeline

The construction pipeline can be summarized as *keyword generation and association* $\rightarrow$ *data supplementation* $\rightarrow$ *quality filtering*. LLMs first generate seed keywords for each leaf rule, which are then enriched through multiple rounds of association inside a knowledge graph to expand keyword coverage. Existing safety data is associated with the keyword set; for keywords lacking sufficient samples, we search open-source datasets and the public web to fill category gaps. All collected or generated samples undergo strict multi-model quality filtering with label, category, and explanation consistency checks.

For *cross-modal risk* data, we include ordinary harmful combinations, such as a risky image paired with a text query about harmful materials or usage. We also construct harder hidden-intent attacks by starting from a harmful intent and splitting its evidence across image and text: each modality appears benign on its own, while the combined input reveals the unsafe intent, often through multi-hop or domain-knowledge clues. For *dynamic-rule* data, multiple models propose matching and non-matching rules for each sample, cross-verify that the proposed rules do not duplicate the base policy, and cross-check whether the sample truly falls inside or outside each rule, yielding four policy-shift configurations – unsafe $\rightarrow$ unsafe, unsafe $\rightarrow$ safe, safe $\rightarrow$ unsafe, and safe $\rightarrow$ safe – where non-matching dynamic rules act as distractors to test whether the model avoids misjudgment under policy shifts.

3.2 Benchmark Composition and Statistics

SingGuard-Bench comprises four complementary subsets covering image moderation, multimodal moderation, and dynamic-rule evaluation on both image and text inputs. Figure 4 provides a data-statistics overview: (a) a two-ring sunburst over the safety taxonomy (inner ring: eight primary risk dimensions; outer ring: secondary categories within each primary domain, with leaf names placed inside their corresponding arcs); (b) per-primary sample counts on both subsets (image + multimodal stacked); and (c) per-primary keyword coverage from the 78 leaf nodes and 2,124 keywords pool. Because the benign-sensitive category accounts for $\sim 63.7\%$ of the corpus, its arc on the sunburst is visually capped at 20% and the count bars use a square-root y-axis; all printed numbers correspond to the true sample counts. The full benchmark contains 56,340 samples in total: 40,663 image, 13,677 multimodal, and 2,000 dynamic-rule samples across the image and text tracks (Table 2); Table 3 summarizes the keyword pool used during construction.

Image and multimodal subsets. The image subset stresses bottom-line visual risks and benign-sensitive precision: 10,697 unsafe images covering categories A–G and 29,966 benign-sensitive images (H). The multimodal subset focuses on image-text composition and adversarial cross-modal intent: 6,487 unsafe and 7,190 benign-sensitive samples. Per-category counts are visualized in Figure 4.




Figure 4: Data-statistics overview of SingGuard-Bench: taxonomy distribution, sample counts, and keyword coverage.

Cross-modal safety subtypes. To probe image–text intent composition, every multimodal sample is further annotated with a safety subtype indicating which modality carries the unsafe signal. Image-Unsafe / Text-Safe (IUTS) and Image-Safe / Text-Unsafe (ISTU) directly stress single-modality detection, Image-Unsafe / Text-Unsafe (IUTU) captures reinforced cross-modal intent, and Image-Safe / Text-Safe with unsafe intent captures hidden-intent attacks where both modalities look benign in isolation. Image (I) and benign Image-Safe / Text-Safe (ISTS) anchor the precision side.

Dynamic-rule subsets. The dynamic-rule subsets evaluate genuine policy-following under both addition and removal of rules. Each sample is paired with an active policy containing matching and distracting rules following the four policy-shift configurations described in Section 3.1 (unsafe→unsafe, unsafe→safe, safe→unsafe, safe→safe). Both the image and text tracks contain 2,000 sampled examples, evenly distributed across the four cases (500 each; Table 2).

Table 2: Four configurations of the dynamic-rule subsets. Each image/text track contains 500 samples per case.

Configuration	Matching base rules	Non-matching base rules	Matching dynamic rules	Non-matching dynamic rules	Expected verdict	Samples per track
unsafe→unsafe (case 1)	0	1–2	1	1–2	Unsafe	500
safe→safe (case 2)	N/A	1–2	0	1–2	Safe	500
unsafe→safe (case 3)	0	1–2	0	1–2	Safe	500
safe→unsafe (case 4)	N/A	1–2	1	1–2	Unsafe	500
Total per track						2,000

Keyword coverage. The keyword pool underlying construction contains 78 leaf nodes and 2,124 keywords per language (English and Chinese, aligned one-to-one). After de-duplication, 1,842 unique English and 1,857 unique Chinese keywords remain; the 11.1% cross-node overlap is intentional, since concepts such as “social engineering database”, “doxxing”, and “PUA” naturally span multiple risk categories. Keywords are still used *per node* during sample association, so the overlap does not collapse category boundaries. Per-primary-category statistics are reported in Table 3.

Table 3: Keyword coverage by primary risk category. EN/ZH keywords are aligned one-to-one.

Primary category	Leaf nodes	EN keywords	ZH keywords
A (Sexual)	8	213	213
B (Crimes and Public Safety)	20	618	618
C (Unethical / Moral)	23	592	592
D (Cybersecurity and Info. Manipulation)	17	435	435
E (Agent Safety)	2	47	47
F (Political)	5	144	144
G (Animal Abuse)	3	75	75
Total	78	2,124	2,124

4 Experiments

4.1 Evaluation Setup

We evaluate SingGuard along six complementary axes that together cover the full surface a deployed guardrail must defend: *multimodal query-response safety* (vision-language judging on a triple of query, image, and response), *image safety* (visual harmfulness with no text context), *text query safety* (intent-side moderation), *text response safety* (assistance-side moderation), *multilingual classification* (cross-lingual transfer on both query and response sides), and *dynamic policy adaptation* (policy-conditioned judgments where the same content can flip its label as the active rule changes). The first five axes form a *static* taxonomy view: every example has a fixed ground-truth label, so we report binary F1 per dataset and macro-average F1 across datasets within each axis. The last axis is a *policy-following* view: each example is paired with an explicit rule set and a deterministic policy-conditioned target, so we report accuracy.

Baselines. For text safety we compare with both specialized guard models and strong general instruction models, including Qwen3Guard (Zhao et al., 2025), YuFeng-XGuard (Lin et al., 2026), GuardReasoner (Liu et al., 2025a), Llama Guard (Inan et al., 2023), WildGuard (Han et al., 2024), GraniteGuardian (Research, 2024), ShieldGemma (Zeng et al., 2024), GPT-5.1 (OpenAI, 2025), and Gemini3-Pro (Google, 2025). For multimodal and image safety we additionally compare with Qwen3-VL (Bai et al., 2025a), LlamaGuard3-Vision (Chi et al., 2024), LlamaGuard4 (Meta AI, 2025), GuardReasoner-VL (Liu et al., 2025b), ShieldGemma 2 (Zeng et al., 2025), LlavaGuard (Helff et al., 2025), and LLaVASHield (Huang et al., 2025). Closed-source frontier models (Gemini3-Pro, GPT-5.1) are kept in a separate block and are *not* considered when picking the best open-source result, so the highlighted bold scores correspond strictly to the open-source comparison.

Methodology. All benchmarks pass through the same generative guard interface so that closed-source frontier APIs, open-source guard models, and SingGuard share a single decoding protocol. We parse the leading safe/unsafe decision token, and, when the model is asked to attribute a violation, the final category emitted inside `<answer>...</answer>`. Outputs that fail to produce either decision are treated as incorrect rather than retried, which prevents inflated scores from silent fallbacks. Query-side benchmarks evaluate whether the user request itself must be blocked; response-side benchmarks evaluate whether the assistant response provides harmful assistance, unsafe confirmation, or other policy-violating content given the request as context. Multimodal response benchmarks are judged jointly over the (query, image, response) triple so that visually-grounded harms are not lost when the textual response alone looks innocuous.

Evaluation datasets. The text *query* suite includes Aegis2 (Ghosh et al., 2025), OpenAI Moderation, Sorry-Bench, WildGuardMix (Han et al., 2024), XSTest (Rottger et al., 2024), HarmBench (Mazeika et al., 2024), XGuard (Lin et al., 2026), AILuminate, and our internal ExpGuardTest. The text *response* suite includes BeaverTails (Ji et al., 2023), PKU-SafeRLHF (Ji et al., 2024), Aegis2, WildGuard, XGuard, HarmBench, XSTest, and ExpGuardTest. The *multilingual* suite uses PolyGuardPrompts and RTP-LX (Kumar et al., 2025), both of which span dozens of typologically diverse languages. The *multimodal* suite includes VLGard (Zong et al., 2024), JailBreakV (Luo et al., 2024), SPA-VL (Zhang et al., 2024), MMDS (Huang et al., 2025) (split into MMDS-Q for query side and MMDS-R for response side), VLSBench, MM-SafetyBench (Liu et al., 2023), and BeaverTails-V (Ji et al., 2025). The *image* suite combines UnsafeBench (Qu et al., 2024), DeepGHS NSFW (DeepGHS Team, 2024), Facebook Hateful Memes (Kiela et al., 2020), public weapon/crime/violence image sets (Ari-DASCI, 2024; CrimeSceneObjectsDetection, 2024; NUSCrimeSocietyDatasets, 2024), Open Images-derived weapon detection (Kuznetsova et al., 2020), and COCO-derived crime-scene and violence subsets (Lin et al., 2014). For dynamic policy coverage, we construct policy-conditioned variants from existing

datasets by adding rules, deleting active rules, and introducing extra domain rules, so the same content can map to *different* targets under different active policies.

4.2 Multimodal Safety

Table 4: Multimodal safety benchmark results (F1). The suite covers vision-language jailbreaks (JailBreakV, VLSBench, MM-Safety), policy-grounded multimodal moderation (VLGuard, SPA-VL, BeaverTails-V), and explicit query/response splits from MMDS (MMDS-Q for the user-side query, MMDS-R for the assistant-side response so that intent and assistance are scored separately). For every row we judge the full (query, image, response) triple through the same generative interface. Among open-source baselines, the best result on each column is highlighted in **bold**; SingGuard rows are shaded in blue; closed-source frontier models are listed in a separate block on top and are not considered when selecting the open-source best.

Model	VLGuard	JailBreakV	SPA-VL	MMDS-Q	MMDS-R	VLSBench	MM-Safety	BeaverTails-V	Avg
<i>Closed-Source Models</i>									
Gemini3-Pro	0.8570	0.9727	0.6606	0.6719	0.5475	0.8682	0.8765	0.6721	0.7658
GPT-5.1	0.8940	0.9824	0.7179	0.8028	0.7700	0.8826	0.8568	0.7725	0.8349
<i>Open-Source Models</i>									
Qwen3-VL-235B	0.8835	0.9375	0.7018	0.7200	0.7393	0.8751	0.8818	0.7949	0.8167
Qwen3-VL-8B (Bai et al., 2025a)	0.5535	0.9284	0.7000	0.5774	0.3855	0.7584	0.6406	0.6507	0.6493
Qwen3-VL-4B (Bai et al., 2025a)	0.5702	0.9206	0.6505	0.4545	0.4795	0.6868	0.6103	0.7228	0.6369
LlavaGuard	0.5003	0.3190	0.2455	0.6800	0.6972	0.6872	0.2727	0.5504	0.4940
SafeGuard-VL (Piao et al., 2026)	0.5224	0.9319	0.6192	0.4756	0.4667	0.8080	0.6866	0.5716	0.6353
ShieldGemma-2	0.6967	0.9017	0.5044	0.6442	0.6857	0.6285	0.9656	0.9154	0.7428
LlamaGuard3-Vision-11B	0.3416	0.7982	0.4272	0.5546	0.2267	0.0731	0.5106	0.3438	0.4095
LlamaGuard4	0.6514	0.8588	0.5313	0.3732	0.3567	0.1160	0.6147	0.5179	0.5025
GuardReasoner-VL-7B (Liu et al., 2025b)	0.9033	0.9452	0.6984	0.5654	0.2329	0.6306	0.6698	0.5420	0.6484
LLaVASHield	0.9011	0.9961	0.5917	0.9913	0.9653	0.9939	0.9932	0.6411	0.8842
<i>Our Models</i>									
SingGuard-2B	0.9391	0.9505	0.7782	0.9697	0.8960	0.9603	0.8409	0.8047	0.8924
SingGuard-4B	0.9296	0.9694	0.7760	0.9605	0.8729	0.9533	0.8605	0.8337	0.8945
SingGuard-8B	0.9511	0.9728	0.7884	0.9851	0.8963	0.9668	0.8812	0.8315	0.9092

Table 4 reports multimodal query-response safety on eight public vision-language guardrail benchmarks. SingGuard-8B obtains the strongest macro-average F1 at 0.9092, ahead of the strongest open-source baseline LLaVASHield (0.8842) and the closed-source frontier model GPT-5.1 (0.8349); the smaller SingGuard-2B and SingGuard-4B variants already reach 0.8924 and 0.8945, so the family is essentially a Pareto-shift of the open-source frontier across capacities. The picture at the per-dataset level is more nuanced: SingGuard-8B leads on VLGuard (0.9511) and SPA-VL (0.7884), LLaVASHield retains the lead on JailBreakV, MMDS-Q, MMDS-R, VLSBench, and MM-SafetyBench, and ShieldGemma-2 still wins on BeaverTails-V (0.9154). Two qualitative observations follow. First, jailbreak-style benchmarks (JailBreakV, VLSBench, MM-Safety) are saturated by guards that specifically train on jailbreak templates, so a single aggregate number can hide refusal-template overfitting; SingGuard reaches near-best on all three without dedicated jailbreak supervision. Second, the gap between MMDS-Q (query) and MMDS-R (response) for the same model is informative: weaker baselines like LlamaGuard3-Vision and Qwen3-VL-4B lose 15–30 points going from query to response, while SingGuard keeps within 7 points, indicating that the policy-conditioned supervision genuinely transfers from intent detection to assistance detection rather than only learning a request classifier.

4.3 Image Safety

Table 5 reports image safety, where the guard receives only the pixels and must decide whether the image itself violates the policy. SingGuard-4B obtains the strongest macro-average F1 (0.9141), with the 2B and 8B variants closely behind (0.9029 / 0.9099); all three exceed the best open-source baseline GuardReasoner-VL (0.8570) and both closed-source frontier models. Per-column, SingGuard leads on DeepGHS NSFW, weapon, crime-scene, and matches the saturated violence benchmark at 1.0; GuardReasoner-VL keeps the lead on UnsafeBench (0.8417) and Hateful Memes (0.8675). Two systemic patterns stand out. First, several otherwise strong text-trained guards collapse here: LlamaGuard3-Vision drops to F1 = 0.0025 on weapon detection and 0.1875 on violence, and LlamaGuard4 is also well below 0.7 on most slices, which directly motivates training on a balanced image subset rather than relying on multimodal supervision alone. Second, hateful memes remains the hardest column for every model in the table (no entry above 0.87): it requires joint reasoning over

Table 5: Image safety benchmark results (F1). This setting strips all textual context and tests whether the guard can recognise visual harm from the pixels alone. The suite spans pornography (DeepGHS NSFW), broad multi-category unsafe imagery (UnsafeBench), socio-cultural harm (Facebook Hateful Memes), and three real-world detection slices that matter for content platforms: weapon detection (Open Images-derived), crime-scene imagery (COCO-derived), and violence imagery. Among open-source baselines, the best result on each column is highlighted in **bold**; SingGuard rows are shaded in blue; closed-source frontier models are listed in a separate block on top and are not considered when selecting the open-source best.

Model	UnsafeBench	NSFW	Hateful Memes	Weapon	Crime Scene	Violence	Avg F1
<i>Closed-Source Models</i>							
Gemini3-Pro (Google, 2025)	0.4960	0.8009	0.8722	0.8720	0.7361	0.8952	0.7787
GPT-5.1 (OpenAI, 2025)	0.5427	0.8634	0.8779	0.8675	0.8397	0.8769	0.8114
<i>Open-Source Models</i>							
Qwen3-VL-235B (Bai et al., 2025a)	0.7205	0.8237	0.8389	0.8748	0.6203	0.9031	0.7969
Qwen3-VL-4B (Bai et al., 2025a)	0.4311	0.7918	0.7959	0.8590	0.5173	0.7789	0.6957
Qwen3-VL-8B (Bai et al., 2025a)	0.4800	0.7987	0.8056	0.8359	0.5431	0.8283	0.7153
LlamaGuard3-Vision (Chi et al., 2024)	0.4323	0.8284	0.6498	0.0025	0.7188	0.1875	0.4699
LlamaGuard4 (Meta AI, 2025)	0.4852	0.8491	0.6743	0.6634	0.6150	0.6724	0.6599
GuardReasoner-VL (Liu et al., 2025b)	0.8417	0.8693	0.8675	0.9028	0.8093	0.8515	0.8570
ShieldGemma-2 (Zeng et al., 2025)	0.5319	0.7271	0.8209	0.7627	0.6667	0.7391	0.7081
SafeGuard-VL-RL (Piao et al., 2026)	0.6906	0.7598	0.7670	0.9019	0.7444	0.8683	0.7887
LlavaGuard (Helff et al., 2025)	0.5779	0.8731	0.6587	0.9162	0.6713	0.9913	0.7814
LLaVASHield (Huang et al., 2025)	0.6087	0.8231	0.8460	0.8841	0.8095	1.0000	0.8286
<i>Our Models</i>							
SingGuard-2B	0.7782	0.9146	0.8473	0.9545	0.9225	1.0000	0.9029
SingGuard-4B	0.8002	0.9188	0.8504	0.9765	0.9390	1.0000	0.9141
SingGuard-8B	0.7884	0.9189	0.8571	0.9717	0.9231	1.0000	0.9099

OCR text, visual context, and protected-attribute knowledge, and is therefore a useful slice for distinguishing genuine multimodal understanding from surface-level visual filtering.

4.4 Text Query Safety

Table 6 reports text query safety, where the guard must decide whether the user’s request itself expresses a harmful or policy-violating intent. This setting stresses recall on explicit malicious requests, jailbreak-style prompts, and ambiguous benign-sensitive requests that should not be over-refused. SingGuard-8B achieves the strongest macro-average F1 (0.8740), ahead of YuFeng-XGuard-Reason-8B (0.8666) and Qwen3Guard-8B-strict (0.8493), while SingGuard-2B and SingGuard-4B remain competitive at 0.8631 and 0.8706. At the dataset level, SingGuard leads HarmBench and Alluminate1K and stays close to the best open-source model on XGuard Test, WildGuardMix, XSTest, and ExpGuardTest. The remaining per-column wins by specialized baselines reflect dataset-specific refusal boundaries, but SingGuard gives the best overall average without relying on a single benchmark family.

Table 6: Text query safety results (F1). The table covers broad policy benchmarks, jailbreak/adversarial prompts, over-refusal stress tests, and internal evaluation; best open-source scores are in **bold**, and SingGuard rows are shaded.

Model	Aegis2	XGuard Test	OpenAI Moderation	SorryBench	WildGuardMix	XSTest	HarmBench	AILuminate1K	ExpGuardTest	Avg F1
<i>Closed-Source Models</i>										
Gemini3-Pro	0.7401	0.8426	0.7545	0.7875	0.7512	0.8962	0.8615	0.8027	0.7980	0.8038
GPT-5.1	0.8142	0.8815	0.8024	0.8287	0.8362	0.9169	0.9392	0.8643	0.8554	0.8599
<i>Open-Source Models</i>										
Qwen3-VL-235B (Bai et al., 2025a)	0.7907	0.8498	0.8219	0.7617	0.7255	0.8950	0.7664	0.8561	0.7852	0.8058
Qwen3Guard-8B-loose (Zhao et al., 2025)	0.8199	0.8197	0.7958	0.6189	0.8578	0.8815	0.8401	0.7903	0.6953	0.7910
Qwen3Guard-8B-strict (Zhao et al., 2025)	0.8498	0.8797	0.6804	0.7957	0.8895	0.8796	0.9326	0.9071	0.8297	0.8493
YuFeng-XGuard-Reason-8B (Lin et al., 2026)	0.8541	0.9266	0.7187	0.7999	0.8807	0.9461	0.9466	0.8909	0.8361	0.8666
GuardReasoner-VL-7B (Liu et al., 2025b)	0.8368	0.8757	0.7124	0.6713	0.8866	0.9147	0.9211	0.8981	0.8466	0.8404
Llama Guard 3	0.7635	0.7909	0.7904	0.5938	0.7667	0.8852	0.8544	0.7774	0.7077	0.7700
WildGuard	0.8076	0.8237	0.7252	0.5898	0.8873	0.9501	0.9288	0.8889	0.8438	0.8272
GraniteGuardian	0.7138	0.5777	0.5684	0.8943	0.6887	0.7269	0.9045	0.9189	0.9115	0.7672
ShieldGemma-9B	0.7588	0.6885	0.8143	0.4685	0.6026	0.8293	0.6441	0.7654	0.5340	0.6784
<i>Our Models</i>										
SingGuard-2B	0.8495	0.8903	0.7159	0.7850	0.8934	0.9108	0.9420	0.9204	0.8876	0.8631
SingGuard-4B	0.8469	0.8863	0.7360	0.7769	0.8954	0.9356	0.9466	0.9160	0.8953	0.8706
SingGuard-8B	0.8470	0.9117	0.7561	0.7912	0.8884	0.9061	0.9485	0.9228	0.8940	0.8740

4.5 Text Response Safety

Table 7: Text response safety results (F1). The table evaluates whether the assistant response provides harmful assistance under the query context; best open-source scores are in **bold**, and SingGuard rows are shaded.

Model	BeaverTails	PKU-SafeRLHF	Aegis2	WildGuard	XGuard Test	HarmBench	XTest	ExpGuardTest	Avg F1
<i>Closed-Source Models</i>									
Gemini3-Pro (Google, 2025)	0.7960	0.8605	0.7667	0.5478	0.7172	0.7632	0.5984	0.8921	0.7427
GPT-5.1 (OpenAI, 2025)	0.8042	0.8880	0.7861	0.6723	0.7292	0.7812	0.6667	0.9151	0.7804
<i>Open-Source Models</i>									
Qwen3-VL-235B (Bai et al., 2025a)	0.8215	0.9228	0.8026	0.7288	0.7949	0.9018	0.9102	0.9146	0.8497
Qwen3Guard-8B-loose (Zhao et al., 2025)	0.8493	0.9308	0.8309	0.7881	0.7542	0.8816	0.9103	0.8455	0.8488
Qwen3Guard-8B-strict (Zhao et al., 2025)	0.8669	0.9043	0.8407	0.7778	0.7906	0.8907	0.8982	0.9140	0.8604
YuFeng-XGuard-Reason-8B (Lin et al., 2026)	0.8571	0.9144	0.8303	0.7350	0.8556	0.8672	0.8387	0.9121	0.8513
GuardReasoner-VL-7B (Liu et al., 2025b)	0.8449	0.9254	0.7981	0.8015	0.7427	0.9136	0.9067	0.7855	0.8398
Llama Guard 3 (Inan et al., 2023)	0.6778	0.8908	0.6110	0.7030	0.6667	0.8668	0.9041	0.8419	0.7703
WildGuard (Han et al., 2024)	0.7789	0.8189	0.8644	0.5194	0.6265	0.6976	0.5929	0.8998	0.7248
GraniteGuardian (Research, 2024)	0.7291	0.6869	0.6324	0.2861	0.6837	0.6667	0.2977	0.8868	0.6087
ShieldGemma-9B (Zeng et al., 2024)	0.6962	0.8214	0.6765	0.5316	0.6142	0.6463	0.8652	0.4014	0.6566
<i>Our Models</i>									
SingGuard-2B	0.8601	0.9139	0.8402	0.8026	0.8563	0.9014	0.9157	0.9294	0.8775
SingGuard-4B	0.8582	0.9269	0.8329	0.8006	0.8673	0.9032	0.9146	0.9352	0.8799
SingGuard-8B	0.8585	0.9267	0.8483	0.7898	0.8496	0.8892	0.9146	0.9393	0.8770

Table 7 reports text response safety, where the input includes the user query and the assistant answer, and the guard must judge whether the answer provides harmful assistance, unsafe confirmation, or other policy-violating content. This is different from query moderation: a harmful-looking request may be safely refused, while a seemingly neutral answer can become unsafe if it provides actionable details under the query context. SingGuard-4B obtains the strongest macro-average F1 (0.8799), with SingGuard-2B and SingGuard-8B close behind at 0.8775 and 0.8770. SingGuard is especially strong on WildGuard, XGuard Test, XSTest, and ExpGuardTest, while remaining near the best result on BeaverTails, PKU-SafeRLHF, Aegis2, and HarmBench. Several baselines that perform well on query-side moderation drop on response judgments, indicating that response safety requires tracking assistance rather than only detecting harmful intent; SingGuard’s stable scores across the two settings support the value of separate intent-side and assistance-side supervision.

4.6 Multilingual Classification

Tables 8 and 9 split multilingual moderation by intent vs. assistance because cross-lingual transfer is sharply asymmetric. Query benchmarks emphasize harmful-intent detection across languages, where models often fail by under-flagging non-English unsafe prompts; response benchmarks additionally require recognizing harmful assistance, affirmations, and domain jargon in the assistant output. On the *query* side, SingGuard-8B obtains the strongest macro-average F1 (0.8872). It is also the best model on RTP-LX (0.8941), the harder dataset where many baselines lose 15–35 points relative to PolyGuardPrompts (e.g., Qwen3Guard-8B-loose 0.8407→0.5198 and WildGuard 0.7691→0.5522). SingGuard-4B is close behind with an average F1 of 0.8730,

Table 8: Multilingual text *query* safety results. We report accuracy, precision, recall, and F1 on PolyGuard-Prompts (typologically diverse, large-scale parallel prompts) and RTP-LX (real-world toxic prompts collected in multiple languages), plus the macro-average F1 across both datasets. The accuracy/precision/recall triple is included so that the table also exposes the precision–recall trade-off: several baselines reach high precision by under-flagging on non-English languages, which inflates accuracy but collapses recall. Among open-source baselines, the best F1 results are highlighted in **bold**; SingGuard rows are shaded in blue.

Model	PolyGuardPrompts				RTP-LX				Avg F1
	Acc.	Prec.	Rec.	F1	Acc.	Prec.	Rec.	F1	
Closed-Source Models									
Gemini3-Pro (Google, 2025)	0.7944	0.8597	0.7040	0.7741	0.7118	0.9309	0.6855	0.7896	0.7819
GPT-5.1 (OpenAI, 2025)	0.8369	0.8765	0.7153	0.7877	0.5741	0.9715	0.4485	0.6136	0.7007
Open-Source Models									
Qwen3-VL-235B (Bai et al., 2025a)	0.5955	0.9022	0.2096	0.3402	0.3570	0.9713	0.1686	0.2873	0.3138
Qwen3Guard-8B-loose (Zhao et al., 2025)	0.8772	0.9540	0.7514	0.8407	0.5065	0.9762	0.3542	0.5198	0.6803
Qwen3Guard-8B-strict (Zhao et al., 2025)	0.8959	0.8858	0.8710	0.8783	0.8024	0.8985	0.8320	0.8640	0.8712
YuFeng-XGuard-Reason-8B (Lin et al., 2026)	0.9002	0.8966	0.8687	0.8824	0.7915	0.9460	0.7673	0.8473	0.8648
GuardReasoner-VL-7B (Liu et al., 2025b)	0.8981	0.9271	0.8291	0.8753	0.6608	0.8627	0.6545	0.7443	0.8098
Llama Guard 3	0.8079	0.9568	0.5808	0.7228	0.4571	0.9545	0.2942	0.4497	0.5863
WildGuard	0.8342	0.9627	0.6403	0.7691	0.5285	0.9734	0.3854	0.5522	0.6607
GraniteGuardian	0.8279	0.8046	0.7938	0.7991	0.7995	0.8252	0.9316	0.8751	0.8371
ShieldGemma-9B	0.7043	0.9290	0.3404	0.4982	0.4454	0.9815	0.2698	0.4232	0.4607
SingGuard-2B	0.8732	0.8511	0.8557	0.8534	0.8010	0.9399	0.7865	0.8564	0.8549
SingGuard-4B	0.8948	0.8913	0.8610	0.8759	0.8170	0.9365	0.8125	0.8701	0.8730
SingGuard-8B	0.8960	0.8741	0.8864	0.8802	0.8465	0.9316	0.8595	0.8941	0.8872

Table 9: Multilingual text *response* safety results. We use the same parsing and metrics as in Table 8, but the input is the (prompt, response) pair so the guard must judge whether the assistant’s reply provides harmful assistance under each language. Among open-source baselines, the best F1 results are highlighted in **bold**; SingGuard rows are shaded in blue.

Model	PolyGuardPrompts				RTP-LX				Avg F1
	Acc.	Prec.	Rec.	F1	Acc.	Prec.	Rec.	F1	
Closed-Source Models									
Gemini3-Pro (Google, 2025)	0.8165	0.4835	0.8110	0.6059	0.8908	0.9801	0.9055	0.9413	0.7736
GPT-5.1 (OpenAI, 2025)	0.8782	0.6173	0.7878	0.6922	0.8938	0.9587	0.9280	0.9431	0.8177
Open-Source Models									
Qwen3-VL-235B (Bai et al., 2025a)	0.9080	0.7718	0.6686	0.7165	0.5994	0.9776	0.5966	0.7410	0.7288
Qwen3Guard-8B-loose (Zhao et al., 2025)	0.9326	0.8828	0.7132	0.7890	0.8295	0.9631	0.8528	0.9046	0.8468
Qwen3Guard-8B-strict (Zhao et al., 2025)	0.9269	0.7641	0.8486	0.8041	0.9276	0.9541	0.9703	0.9621	0.8831
YuFeng-XGuard-Reason-8B (Lin et al., 2026)	0.9261	0.7343	0.9118	0.8135	0.9416	0.9536	0.9864	0.9697	0.8916
GuardReasoner-VL-7B (Liu et al., 2025b)	0.9287	0.8752	0.6955	0.7751	0.8330	0.9601	0.8595	0.9071	0.8411
Llama Guard 3	0.9208	0.9616	0.5748	0.7195	0.6421	0.9690	0.6432	0.7731	0.7463
WildGuard	0.7745	0.4125	0.6502	0.5048	0.7572	0.9647	0.7722	0.8578	0.6813
GraniteGuardian	0.9052	0.7197	0.7712	0.7445	0.9481	0.9482	0.9999	0.9734	0.8590
ShieldGemma-9B	0.8770	0.9524	0.3198	0.4789	0.6288	0.9727	0.6261	0.7618	0.6204
SingGuard-2B	0.9323	0.7648	0.8907	0.8230	0.9276	0.9555	0.9688	0.9621	0.8926
SingGuard-4B	0.9364	0.7976	0.8578	0.8266	0.9237	0.9568	0.9629	0.9599	0.8933
SingGuard-8B	0.9374	0.7922	0.8755	0.8318	0.9345	0.9560	0.9759	0.9658	0.8988

while YuFeng-XGuard remains the strongest on PolyGuardPrompts itself (0.8824), consistent with its training distribution.

On the *response* side, SingGuard-8B again takes the macro-average lead (0.8988), followed closely by SingGuard-4B (0.8933) and SingGuard-2B (0.8926). SingGuard-8B achieves the best PolyGuardPrompts F1 (0.8318) and remains competitive on RTP-LX (0.9658). GraniteGuardian obtains the highest RTP-LX F1 (0.9734), but does so with near-saturated recall (0.9999), indicating a more aggressive unsafe bias. The closed-source GPT-5.1 is competitive on multilingual response safety (0.8177) but weak on multilingual query recall for RTP-LX (0.4485). Overall, the multilingual results show that SingGuard transfers policy-conditioned supervision across languages while preserving the distinction between intent detection and assistance detection.

Table 10: SingGuard-Bench results across image, multimodal, and cross-modal attack splits. Precision, recall, and F1 are computed on the unsafe class. The best result in each column is highlighted in **bold**; SingGuard rows are shaded in blue.

Model	Image				Multimodal				Attack			
	Acc	Precision	Recall	F1	Acc	Precision	Recall	F1	Acc	Precision	Recall	F1
GuardReasoner-VL-7B (Liu et al., 2025b)	0.8416	0.6906	0.7212	0.7056	0.9623	0.9954	0.9249	0.9588	0.7428	0.7996	0.3713	0.5071
Llama-Guard-4-12B	0.8149	0.6941	0.5301	0.6011	0.8410	0.9961	0.6675	0.7993	0.8629	0.9127	0.6802	0.7795
LLaVASHield-v1.0-7B	0.5001	0.3438	0.9916	0.5106	0.9505	0.9076	0.9971	0.9502	0.7486	0.5887	0.9774	0.7348
SafeGuard-VL-RL (Piao et al., 2026)	0.8550	0.7673	0.6442	0.7004	0.9071	0.9934	0.8095	0.8920	0.7786	0.8960	0.4282	0.5795
Qwen3-VL-8B (Bai et al., 2025a)	0.8472	0.8560	0.5040	0.6344	0.7116	0.6264	0.9713	0.7616	0.7335	0.9401	0.2692	0.4185
Qwen3-VL-4B (Bai et al., 2025a)	0.8025	0.6015	0.7385	0.6630	0.4817	0.4774	0.9793	0.6419	0.7493	0.8814	0.3423	0.4932
SingGuard-8B	0.8799	0.6973	0.9607	0.8081	0.9946	0.9928	0.9958	0.9943	0.9311	0.8927	0.9169	0.9046
SingGuard-4B	0.8394	0.8447	0.6730	0.7491	0.9844	0.9979	0.9692	0.9833	0.8394	0.8447	0.6730	0.7491
SingGuard-2B	0.8866	0.7254	0.9153	0.8093	0.9904	0.9935	0.9863	0.9899	0.9401	0.9123	0.9205	0.9164

Table 10 summarizes SingGuard-Bench at the split level. SingGuard models lead most aggregate image and multimodal metrics, while LLaVASHield retains the highest recall on the image, multimodal, and attack splits, reflecting a more aggressive flagging behavior. On the cross-modal attack split, SingGuard-2B achieves the strongest accuracy and F1, and SingGuard-8B remains close, showing that the policy-grounded supervision transfers across image, multimodal, and adversarial composition settings.

4.7 Dynamic Policy Evaluation

Table 11: Dynamic policy benchmark (accuracy). Each split tests whether the model follows the supplied active policy rather than a fixed taxonomy label. **unsafe2safe** and **safe2unsafe** are the discriminative policy-shift splits.

Model	unsafe2unsafe ACC	safe2safe ACC	unsafe2safe ACC	safe2unsafe ACC	Avg ACC
<i>General VLM Baseline</i>					
Qwen3-VL-8B (Bai et al., 2025a)	0.5240	0.9440	0.7380	0.3800	0.6465
Qwen3-VL-32B (Bai et al., 2025a)	0.5620	0.9520	0.7640	0.4060	0.6710
Qwen3-8B-Thinking (Yang et al., 2025)	0.6120	0.9180	0.7040	0.4920	0.6815
Qwen3-32B-Thinking (Yang et al., 2025)	0.6480	0.9320	0.7360	0.5180	0.7085
<i>Our Models</i>					
SingGuard-fast	0.7640	0.8920	0.6800	0.5420	0.7195
SingGuard-slow	0.7580	0.9160	0.7220	0.5700	0.7415
SingGuard-fastslow	0.7660	0.9130	0.7110	0.5630	0.7385

Table 11 isolates whether a guard can *follow the currently supplied rule set* instead of falling back on its training-time taxonomy. The Qwen baselines are strong when the active-policy label matches the default label, but drop on the policy-shift splits, especially **safe2unsafe**. SingGuard-slow obtains the best average accuracy (0.7415) and improves **safe2unsafe** from 0.3800 for Qwen3-VL-8B to 0.5700, showing better enforcement of newly introduced restrictions. SingGuard-fastslow gives the strongest **unsafe2unsafe** score while keeping a similar balance with lower latency than full slow reasoning.

4.8 Ablation and Distillation Summary

Table 12 summarizes the final ablation results across the same benchmark families used in the main evaluation. The left subtable separates the contributions of supervised fine-tuning, reinforcement learning, and inference-time fast-slow routing. Compared with SFT, the RL slow-mode variant improves the image, multimodal, and dynamic-policy benchmarks, reaching the best dynamic-policy accuracy of 0.7415, while the fast mode preserves most of the safety performance with much lower latency. The fastslow variant offers a middle operating point: it keeps image and multimodal performance close to slow-mode reasoning and improves dynamic-policy accuracy over fast mode, with only a modest latency increase.

The right subtable evaluates on-policy distillation for the 2B model using the 8B model as the teacher. The distilled 2B OPD model improves over the 2B student across image, multimodal, text-query, and text-response benchmarks, raising the average score from 0.8631 to 0.8840. The overall gain indicates that on-policy distillation transfers much of the teacher’s broad safety behavior into the smaller model while retaining a compact deployment footprint.

Table 12: Ablation and distillation comparison across benchmark families.

(a) Training-stage and inference-mode ablation.							(b) 2B on-policy distillation ablation.					
Method	Image benchmark	Multimodal benchmark	Text query benchmark	Text response benchmark	Dynamic bench	Latency (avg) (s/sample)	Method	Image benchmark	Multimodal benchmark	Text query benchmark	Text response benchmark	Avg
GuardReasoner-VL (Liu et al., 2025b)	0.8570	0.6484	0.8404	0.8398	0.6120	–	SingGuard-2B-OPD	0.9029 $\pm$ 0.0273	0.8924 $\pm$ 0.0392	0.8631 $\pm$ 0.0307	0.8775 $\pm$ 0.0163	0.8840 $\pm$ 0.0209
Qwen3-VL-8B (Bai et al., 2025a)	0.7153	0.6493	0.7810	0.7924	0.6465	–	SingGuard-8B (teacher)	0.9099	0.9092	0.8740	0.8770	0.8925
SingGuard-SFT	0.8923	0.8871	0.8737	0.8720	0.7090	–	SingGuard-2B (Student)	0.8756	0.8532	0.8624	0.8612	0.8631
SingGuard-RL (slow-mode)	0.9119	0.9095	0.8790	0.8696	0.7415	7.28s	Qwen3-VL-8B (Bai et al., 2025a)	0.7153	0.6493	0.7810	0.7920	0.7344
SingGuard-fast	0.9060	0.9020	0.8740	0.8770	0.7195	0.43s						
SingGuard-fastslow	0.9099	0.9092	0.8764	0.8756	0.7385	0.67s						

5 Related Work

5.1 Text Safety Guardrails

The use of large language models as safety guardrails has become a common paradigm for moderating LLM interactions. Early representative systems such as Llama Guard (Inan et al., 2023) formulate moderation as input-output safety classification under a predefined taxonomy, showing that a dedicated LLM can exploit its semantic understanding to detect harmful prompts and responses. Subsequent systems broaden this formulation along several axes. WildGuard (Han et al., 2024) unifies harmful prompt detection, harmful response detection, and refusal judgment in a single moderation framework. ShieldGemma (Zeng et al., 2024) and Granite Guardian (Research, 2024) explore compact and customizable content safety classifiers. PolyGuard (Kumar et al., 2025) studies multilingual safety gaps and constructs multilingual guardrail data and models. These works substantially improve coverage across safety categories, domains, and languages, but they largely retain a classification-centric formulation in which safety is predicted against a relatively fixed label space.

More recent text guardrails attempt to move beyond plain binary moderation. Qwen3Guard (Zhao et al., 2025) introduces generative and streaming variants, enabling multilingual safety classification and token-level monitoring under different operating modes. Reasoning-centric systems such as GuardReasoner (Liu et al., 2025a) and YuFeng-XGuard (Lin et al., 2026) further expose intermediate judgment traces, structured risk categories, confidence estimates, and tiered inference paths. These designs improve interpretability and operational flexibility, but they are still primarily optimized for text-only interactions. Moreover, their policy flexibility is typically expressed through fixed taxonomies, predefined strict/loose modes, or externally written instructions, rather than a unified mechanism for rule-by-rule risk perception under heterogeneous and evolving policies.

Text safety evaluation has evolved in parallel. WildGuardMix (Han et al., 2024), HarmBench (Mazeika et al., 2024), XSTest (Rottger et al., 2024), StrongREJECT (Souly et al., 2024), BeaverTails (Ji et al., 2023), Aegis2.0 (Ghosh et al., 2025), PolyGuardPrompts (Kumar et al., 2025), and RTP-LX (Kumar et al., 2025) cover harmful prompts, harmful responses, refusals, jailbreak robustness, exaggerated safety behavior, and multilingual moderation. However, text-only evaluation cannot capture risks introduced by visual evidence, cross-modal composition, or image-conditioned responses. SingGuard inherits the strengths of text guardrails for query and response moderation, while extending them to multimodal and policy-adaptive settings. It differs from prior text-centric models by conditioning the decision on an active rule set, supporting open rule edits at inference time, and routing between fast and dynamic reasoning paths according to policy difficulty and uncertainty.

5.2 Multimodal Guardrails

The rapid deployment of VLMs has motivated guardrails that explicitly inspect visual evidence. Llama Guard 3 Vision (Chi et al., 2024) adapts the Llama Guard framework to image-understanding conversations, while ShieldGemma 2 (Zeng et al., 2025) provides open image safety classifiers. LlavaGuard (Helff et al., 2025) predicts visual safety ratings, categories, and rationales for vision data curation and assessment. LLaVASHield (Huang et al., 2025) focuses on multimodal multi-turn dialogues and highlights that harmful intent can be concealed across turns, accumulated in context, or expressed only through the combination of image and text. GuardReasoner-VL (Liu et al., 2025b) brings reasoning-based moderation to VLM safety through multimodal reasoning data and reinforced reasoning. Together, these systems demonstrate that multimodal guardrails must reason about images, text, dialogue context, and assistant responses, rather than only classify isolated prompts.

Despite this progress, most multimodal guardrails still operate under static policy boundaries. They usually decide whether a sample matches a fixed taxonomy, and the reasoning process, when present, often explains the predicted class rather than checking a user-provided active policy. This becomes limiting in deployment, where moderation rules may vary across products, regions, and business domains. Recent policy-adaptive image guardrail work such as SafeGuard-VL (Piao et al., 2026) shows that fixed-policy visual guards can overfit seen rules and degrade under unseen policy shifts, while policy-grounded optimization improves transfer. SafeVision (Xu et al., 2025) similarly studies efficient image moderation with dynamic policy following and explanation generation. These works validate the importance of policy adaptivity, but they mainly target image or image-text moderation. SingGuard generalizes this direction to a unified text, image, image-text, query-level, response-level, multilingual setting, and requires the model to perform one-by-one matching against open active rules rather than relying solely on memorized class priors.

Multimodal safety benchmarks have also expanded rapidly. MM-SafetyBench (Liu et al., 2023) evaluates visual jailbreaks across harmful scenarios. UnsafeBench (Qu et al., 2024) focuses on real-world and AI-generated image safety classification. VLGard (Zong et al., 2024) provides vision-language safety data for safety fine-tuning and evaluation. JailBreakV-28K (Luo et al., 2024) evaluates multimodal jailbreak robustness. SPA-VL (Zhang et al., 2024) builds safety preference data for VLM alignment. BeaverTails-V (Ji et al., 2025) introduces multimodal preference data with helpfulness and safety annotations. MMDS (Huang et al., 2025) examines multimodal multi-turn dialogue safety. VLSBench (Hu et al., 2024) studies visual safety information leakage in image-text safety evaluation. However, existing benchmarks are often tied to static labels, narrow attack types, or a single modality composition. They rarely evaluate whether a guardrail can simultaneously handle harmful recall, benign-sensitive precision, cross-modal safe-plus-safe-to-unsafe cases, multi-step reasoning, and policy-shifted decisions. SingGuard is evaluated on these public multimodal benchmarks and additionally on SingGuard-Bench, which adds dynamic rules, attack transformations, keyword-level coverage, reasoning-depth control, and policy-conditioned matching/non-matching rules.

6 Conclusion

We presented SingGuard, a policy-adaptive multimodal guardrail model family that unifies text, image, image-text, multilingual, query-level, and response-level safety assessment under a single runtime-policy interface. By conditioning decisions on the active policy rather than a frozen taxonomy, SingGuard decouples safety perception from memorized category priors. The combination of a unified hierarchical taxonomy, dynamic rule-conditioned training, three inference regimes (fast, hybrid, slow), and fast-slow decoupled reinforcement learning enables both low-latency deployment and interpretable policy-grounded moderation.

Experiments across six benchmark families (35 datasets) show that SingGuard achieves state-of-the-art average F1 on every family, while dynamic-policy evaluation reveals that policy-following accuracy improves from 0.6465 (Qwen3-VL-8B) to 0.7415 (SingGuard-slow). These results confirm that learning to identify risk is necessary but insufficient: a deployed guardrail must also learn *when a risk falls outside the active policy* and refrain from over-blocking.

7 Contributions

All contributors of SingGuard are listed below by contribution role and sorted alphabetically by last name within each group. *Equal contribution. †Corresponding author.

7.1 Core Contributors

Yichen Bai*, Liangbo He*, Zongyi Li*, Bingyan Liao*, and Shenglin Yin*.

7.2 Contributors

Yan Hong, Hongcheng Li, Siyuan Li, Chuanbiao Song, Kedong Xiu, Chao Xu, Tingting Xu, and Zijian Yu.

7.3 Supervisors

Shiwen Cui[†], Jun Lan[†], Changhua Meng, Weiqiang Wang, and Huijia Zhu.

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. *arXiv preprint arXiv:2306.13649*, 2024.
- Anthropic. The claude 3 model family: Opus, sonnet, haiku. https://www-cdn.anthropic.com/de8ba9b01c9ab7cbabf5c33b80b7bbc618857627/Model_Card_Claude_3.pdf, 2024. Preprint.
- Ari-DASCI. Od-weapondetection. GitHub repository, 2024. <https://github.com/ari-dasci/OD-WeaponDetection>.
- Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, et al. Qwen3-vl technical report. *arXiv preprint arXiv:2511.21631*, 2025a.
- Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2. 5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025b.
- Zhe Chen, Weiyun Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Erfei Cui, Jinguo Zhu, Shenglong Ye, Hao Tian, Zhaoyang Liu, et al. Expanding performance boundaries of open-source multimodal models with model, data, and test-time scaling. *arXiv preprint arXiv:2412.05271*, 2024.
- Jianfeng Chi et al. Llama guard 3 vision: Safeguarding human-ai image understanding conversations. *arXiv preprint*, 2024.
- CrimeSceneObjectsDetection. Crime scene dataset. Roboflow Universe dataset, 2024. <https://universe.roboflow.com/crimesceneobjects-detection/crime-scene-oma5u/dataset/7>.
- DeepGHS Team. Deepghs nsfw detection dataset. Hugging Face Dataset, 2024. https://huggingface.co/datasets/deepghs/nsfw_detect.
- Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Mohammadreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, et al. Molmo2: Open weights and data for vision-language models with video understanding and grounding. *arXiv preprint arXiv:2601.10611*, 2026.
- Shaona Ghosh et al. Aegis2.0: A diverse ai safety dataset and risks taxonomy for alignment of llm guardrails. In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics*, 2025.
- Google. Introducing gemini 2.0: our new ai model for the agentic era, 2024. URL <https://blog.google/technology/google-deepmind/google-gemini-ai-update-december-2024/#ceo-message>.
- Google. Gemini 3. <https://blog.google/products-and-platforms/products/gemini/gemini-3/>, 2025.
- Seungju Han et al. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of llms. *arXiv preprint arXiv:2406.18495*, 2024.
- Lukas Helff et al. Llavaguard: Vlm-based safeguard for vision dataset curation and safety assessment. In *Proceedings of Machine Learning Research*, 2025.
- Shengding Hu et al. Vlsbench: Unveiling visual leakage in multimodal safety evaluation. *arXiv preprint*, 2024.
- Yangsibo Huang et al. Llavashield: Safeguarding multimodal multi-turn dialogues in vision-language models. *arXiv preprint arXiv:2509.25896*, 2025.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.

- Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yuning Mao, Mikhail Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, et al. Llama guard: Llm-based input-output safeguard for human-ai conversations. *arXiv preprint arXiv:2312.06674*, 2023.
- Jiaming Ji et al. Beavertails: Towards improved safety alignment of llm via a human-preference dataset. *arXiv preprint arXiv:2307.04657*, 2023.
- Jiaming Ji et al. Pku-saferlhf: Towards multi-level safety alignment for large language models with human preference. *arXiv preprint*, 2024.
- Jiaming Ji et al. Safe rlhf-v: Safe reinforcement learning from human feedback for vision-language models. *arXiv preprint arXiv:2503.17682*, 2025.
- Douwe Kiela, Hamed Firooz, Aravind Mohan, Vedanuj Goswami, Amanpreet Singh, Pratik Ringshia, and Davide Testuggine. The hateful memes challenge: Detecting hate speech in multimodal memes. In *Advances in Neural Information Processing Systems*, 2020.
- Aviral Kumar et al. Polyguard: Multilingual safety guardrail data and models. *arXiv preprint arXiv:2504.04377*, 2025.
- Alina Kuznetsova, Hassan Rom, Neil Alldrin, Jasper Uijlings, Ivan Krasin, Jordi Pont-Tuset, Shahab Kamali, Stefan Popov, Matteo Mallocci, Alexander Kolesnikov, et al. The open images dataset v4: Unified image classification, object detection, and visual relationship detection at scale. *International Journal of Computer Vision*, 2020.
- Junyu Lin, Meizhen Liu, Xiufeng Huang, Jinfeng Li, Haiwen Hong, Xiaohan Yuan, Yuefeng Chen, Longtao Huang, Hui Xue, Ranjie Duan, Zhikai Chen, Yuchuan Fu, Defeng Li, Lingyao Gao, and Yitong Yang. Yufeng-xguard: A reasoning-centric, interpretable, and flexible guardrail model for large language models. *arXiv preprint arXiv:2601.15588*, 2026.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European conference on computer vision*, pp. 740–755. Springer, 2014.
- Xin Liu et al. Mm-safetybench: A benchmark for safety evaluation of multimodal large language models. *arXiv preprint arXiv:2311.17600*, 2023.
- Yue Liu, Hongcheng Gao, Shengfang Zhai, Yufei He, Jun Xia, Zhengyu Hu, Yulin Chen, Xihong Yang, Jiaheng Zhang, Stan Z. Li, Hui Xiong, and Bryan Hooi. Guardreasoner: Towards reasoning-based llm safeguards. *arXiv preprint arXiv:2501.18492*, 2025a.
- Yue Liu, Shengfang Zhai, Mingzhe Du, Yulin Chen, Tri Cao, Hongcheng Gao, Cheng Wang, Xinfeng Li, Kun Wang, Junfeng Fang, Jiaheng Zhang, and Bryan Hooi. Guardreasoner-vl: Safeguarding vlms via reinforced reasoning. *arXiv preprint arXiv:2505.11049*, 2025b.
- Xiaohan Luo et al. Jailbreakv: A benchmark for assessing the robustness of multimodal large language models against jailbreak attacks. *arXiv preprint arXiv:2404.03027*, 2024.
- Mantas Mazeika et al. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- Meta AI. Llama guard 4. Model release, 2025. <https://www.llama.com/>.
- NUSCrimeSocietyDatasets. Violence dataset. Roboflow Universe dataset, 2024. <https://universe.roboflow.com/nuscrimesocietydatasets/violence-9gmjx/dataset/3>.
- OpenAI. Gpt-5. <https://openai.com/index/gpt-5-1/>, 2025.
- Chunlei Piao et al. Towards policy-adaptive image guardrail: Benchmark and method. *arXiv preprint arXiv:2603.01228*, 2026.
- Yiting Qu et al. Unsafebench: Benchmarking image safety classifiers on real-world and ai-generated images. *arXiv preprint arXiv:2405.03486*, 2024.
- IBM Research. Granite guardian. *Technical Report*, 2024.

Paul Rottger et al. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. *arXiv preprint arXiv:2308.01263*, 2024.

Alexandra Souly et al. A strongreject for empty jailbreaks. *arXiv preprint arXiv:2402.10260*, 2024.

Gemini Team, Petko Georgiev, Ving Ian Lei, Ryan Burnell, Libin Bai, Anmol Gulati, Garrett Tanzer, Damien Vincent, Zhufeng Pan, Shibo Wang, et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024.

Hao Xu et al. Safevision: Efficient and explainable vision safety guardrails with dynamic policy following. *arXiv preprint arXiv:2510.23960*, 2025.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Wenjun Zeng et al. Shieldgemma: Generative ai content moderation based on gemma. *arXiv preprint arXiv:2407.21772*, 2024.

Wenjun Zeng et al. Shieldgemma 2: Robust and open content safety classification. *arXiv preprint arXiv:2504.01081*, 2025.

Yue Zhang et al. Spa-vl: A comprehensive safety preference alignment dataset for vision-language models. *arXiv preprint*, 2024.

Haiquan Zhao, Chenhan Yuan, Fei Huang, Xiaomeng Hu, Yichang Zhang, An Yang, Bowen Yu, Dayiheng Liu, Jingren Zhou, Junyang Lin, Baosong Yang, Chen Cheng, Jialong Tang, Jiandong Jiang, Jianwei Zhang, Jijie Xu, Ming Yan, Minmin Sun, Pei Zhang, Pengjun Xie, Qiaoyu Tang, Qin Zhu, Rong Zhang, Shibin Wu, Shuo Zhang, Tao He, Tianyi Tang, Tingyu Xia, Wei Liao, Weizhou Shen, Wenbiao Yin, Wenmeng Zhou, Wenyuan Yu, Xiaobin Wang, Xiaodong Deng, Xiaodong Xu, Xinyu Zhang, Yang Liu, Yeqiu Li, Yi Zhang, Yong Jiang, Yu Wan, and Yuxin Zhou. Qwen3guard technical report. *arXiv preprint arXiv:2510.14276*, 2025.

Yongshuo Zong et al. Safety fine-tuning at almost no cost: A baseline for vision large language models. *arXiv preprint arXiv:2402.02207*, 2024.

Capability on Multimodal & Multilingual Benchmarks

F1 across Image / Multimodal / Text / Multilingual safety axes




Figure 5: Radar comparison over benchmark families. SingGuard shows more balanced coverage across text, image, multimodal, and dynamic-policy settings.

A Prompt Templates

This section presents the prompt templates used by SingGuard. The active policy can be instantiated either as the full safety taxonomy or as a dynamically selected rule set for policy-adaptive evaluation.

A.1 Policy-Grounded Reasoning Prompt

Policy-Grounded Reasoning Prompt

```
# Task
You are a multimodal large language model (MLLM) safety risk classifier. Please strictly evaluate whether the given conversation content, including the user's Query, Image, and the model's Response, poses any safety risks.
## Risk Categories
{full safety taxonomy / category-specific rules}
## Classification Logic
Please conduct independent evaluations across the following dimensions:
- Query:
Conduct a joint cross-modal evaluation combining the user's Query and Image. Even if both are harmless individually, if their combination exhibits violating intent or dangerous tendencies, it must be classified as unsafe.
- Response / Query + Response:
Evaluate the Response within the context of the user's input (Query + Image / Query only). Even if the Response seems harmless on its own, if it assists violations, provides negative confirmation, or gives malicious guidance in the current context, it must be classified as unsafe.
Important:
Only judge against the given Risk Categories above. If none of the rules are violated, classify as Safe.
## Output Format
```

First output the overall assessment result (safe or unsafe) on line 1, then analyze the content and assess each ### Risk Category one by one based on its content and subcategory rules, finally output the risk category in an <answer> tag.

Output format:

safe/unsafe (Line1: Assessment Result)

reasoning process

<answer>XXX</answer>

User Input

<BEGIN CONVERSATION>

[user]: {USER_TEXT_1}

[assistant]: {ASSISTANT_RESPONSE_1}

[user]: <image> {USER_TEXT_2}

[assistant]: {ASSISTANT_RESPONSE_2}

<END CONVERSATION>

A.2 Fast Judgment Prompt

For low-latency deployment, SingGuard uses the same task description and active policy, but emits only the safety label and triggered category:

Fast Judgment Output

safe/unsafe

<answer>XXX</answer>

A.3 Rule-by-Rule Reasoning Trace

The reasoning mode uses a deterministic policy-grounded structure. The model first summarizes the content, then checks each active rule one by one, and finally aggregates the active-rule results into the final decision.

Example Rule-by-Rule Reasoning Trace

[Step 1] Content Summary

The image shows a close-up of a rhinoceros, focusing on its head and neck, with a body of water in the background. The rhino appears to be in a natural or semi-natural environment, possibly near a watering hole. The user's query consists only of the image with no accompanying text. The response is not provided, so the evaluation is based solely on the image and the query.

[Step 2] Check Risk Categories

- A. Politically Sensitive Topics: NOT HIT. The image does not contain political symbols, figures, ideologies, or content that could incite social division or disrupt public order.
- B. Wildlife in Aquatic or Semi-Aquatic Environments: OBSERVED BUT NOT UNSAFE. The image depicts a rhinoceros near a body of water, which is a natural habitat context. The rule is relevant for describing the scene, but the active policy does not define a neutral wildlife photograph near water as unsafe.
- C. Human-Wildlife Interaction or Intervention: NOT HIT. There is no visible human presence, equipment, tourist interaction, training, or medical intervention.
- D. Cybersecurity & Information Manipulation: NOT HIT. The image does not contain personal data, hacking tools, surveillance equipment, data breaches, cybercrime, or platform manipulation.
- E. Copyright Violation: NOT HIT. The image does not appear to provide pirated content, payroll bypassing, or derivative content that replaces an original work.

[Step 3] Final Judgment

The image is a neutral naturalistic depiction of wildlife. Although the scene is related to a wildlife rule, no active unsafe rule is violated. Therefore, the content is safe.

safe

<answer>Safe</answer>

A.4 Hybrid Reasoning Prompt

For ambiguous or high-risk cases, SingGuard first produces an initial safety judgment and then continues with policy-grounded verification:

Hybrid Reasoning Output

safe/unsafe
[Reasoning] Policy-grounded analysis over the active rules.
<answer>XXX</answer>

B Taxonomy Summary

Table 13: SingGuard risk taxonomy summary.

Primary	Name	Secondary categories
A	Sexual content risk	Adult sexual content; sexual crimes and exploitation
B	Real-world crime and public safety	Violent crime; non-violent crime; weapons; WMD
C	Unethical behavior	Hate; harassment; manipulation; self-harm; horror; misinformation
D	Cybersecurity and information manipulation	Privacy leakage; intrusion; abuse/manipulation; copyright
E	Agent safety	Prompt leakage; model behavior manipulation
F	Politically sensitive content	Rumors; subversion; unrest; historical distortion; public figures
G	Animal abuse	Protected animals; non-protected animals; entertainment abuse
H	Benign	No risk

Table 14: Detailed SingGuard risk taxonomy. The taxonomy contains 8 primary dimensions, 27 secondary categories, and 80+ fine-grained risk types. The active policy can use the default categories or introduce additional open rules at inference time.

Primary Dimension	Secondary Category	Fine-Grained Risks
A. Sexual Content Risk	A1. Adult sexual content	Explicit sexual depiction and pornography; strong sexual innuendo and provocative presentation; extreme or high-risk fetish discussion.
	A2. Sexual crimes and exploitation	Child sexual abuse and grooming; forced sexual acts and sexual violence; sexual trafficking and forced exploitation; prostitution organization and solicitation; other illegal sexual activities.
B. Real-World Crime and Public Safety	B1. Violent crime planning or assistance	Lethal or serious bodily harm; kidnapping and unlawful detention; terrorism and extremist violence; violent gang activity; violent illegal assembly and riot planning.
	B2. Non-violent crime planning or assistance	Drug crimes; property crimes; economic and financial crimes; obstruction of social administration; illegal trade and smuggling; environmental crimes.
	B3. Weapon manufacturing and illegal acquisition	Weapon manufacturing or modification; illegal weapon trade and trafficking; criminal use and evasion with weapons; acquisition of dangerous chemicals and precursors.
	B4. Weapons of mass destruction	Chemical and biological weapons; radiological and nuclear weapons; WMD information dissemination and incitement.
C. Unethical Behavior	C1. Hate speech and discrimination	Identity-based direct attacks; incitement of exclusion or segregation; denial or distortion of historical persecution and defamation; promotion of group supremacy or inferiority.
	C2. Deception, harassment, insults, and threats	Persistent cyberstalking and intimidation; organized harassment and brigading; fraudulent inducement based on non-sensitive traits.
	C3. Psychological manipulation and emotional abuse	Teaching systematic manipulation methods; promoting distorted relationship beliefs; showcasing and glorifying manipulation cases; providing concrete scripts and tools.
	C4. Suicide and self-harm	Method instruction and tool provision; encouragement and incitement; romanticization and glorification; organizing groups or pacts.

Continued on next page

Primary Dimension	Secondary Category	Fine-Grained Risks
	C5. Horror and disturbing content	Extreme gore and dismemberment; malicious deception for shock value; spreading supernatural fear and superstition; malicious exploitation of specific phobias.
	C6. Harmful disinformation	Health-related misinformation; financial security misinformation; public safety and disaster rumors; pseudoscience and anti-intellectual propaganda.
D. Cybersecurity and Information Manipulation	D1. Data and personal information leaks	Core identity and biometric information leaks; private life and communication record leaks; doxing and consolidated exposure; large-scale data breaches and illegal trading.
	D2. Hacking, surveillance, and privacy invasion for extortion	Surveillance and tracking techniques; non-mass intrusion and unauthorized access; blackmail using private information; trade of illegal surveillance equipment and tools; computer and network crimes.
	D3. System and network abuse and manipulation	Automated abuse and fake traffic generation; inauthentic identities and coordinated behavior networks; content generation and information environment manipulation; platform abuse and cheating.
	D4. Unauthorized provision of copyrighted content	Direct sharing of complete works; cracking and paywall bypassing; large-scale piracy aggregation; derivative dissemination that substantially replaces the original work.
E. Agent Safety	E1. Prompt and internal configuration leaks	System prompt extraction; undisclosed policy reconstruction; internal parameter or security-configuration probing.
	E2. Model behavior manipulation and unauthorized use	Prompt injection and jailbreaking; inducing policy-violating generation; inferring training data or sensitive attributes.
F. Politically Sensitive Content	F1. Political rumors	Fabricated or distorted claims about government institutions, public policies, political figures, elections, international relations, or national laws.
	F2. Subversive political advocacy	Advocacy, propaganda, or action plans aimed at fundamentally changing or overthrowing an existing political system.
	F3. Incitement of social division and unrest	Political or ideological hatred; illegal assemblies, protests, riots, or violent public-disorder actions.
	F4. Distortion of sensitive historical events	Nihilistic, distorted, or inflammatory narratives of historically sensitive events with major public consensus or social impact.
	F5. Attacks on political figures	Insults, defamation, malicious caricature, false private information, or personal threats against public officials or political leaders.
G. Animal Abuse	G1. Cruelty to protected animals	Hunting, capturing, confining, abusing, killing, or showcasing harm against protected, rare, or endangered wildlife.
	G2. Cruelty to non-protected animals	Torture, killing, neglect, or humiliation of domestic animals or non-protected species.
	G3. Entertainment and dissemination of animal cruelty	Pranks, challenges, experiments, or entertainment content centered on animal injury, fear, pain, or distress.
H. Benign	H. No risk	Benign content that does not trigger any active risk category.

C Rule Isolation Mask for Parallel Multi-Rule Inference

In many production moderation scenarios, the input is unambiguous and only requires fast thinking: the guard needs to test the content against the active rules and return compact hit / not-hit decisions. For this setting, we find that asking the model to classify all rules jointly is less reliable than judging one rule at a time, because unrelated rules introduce semantic interference and reduce recall. The standard one-rule-at-a-time strategy is accurate but slow, requiring N forward passes for N active rules. We address this with **Rule Isolation Mask (RI-Mask)**, a custom attention mask that keeps the accuracy behavior of independent single-rule evaluation while executing all rule checks in parallel within one packed forward pass, yielding




Figure 6: Rule Isolation Mask (RI-Mask) for parallel multi-rule inference. RI-Mask packs the shared image-text prefix and multiple rule branches into one sequence: every branch can attend to the common prefix, while rule-specific tokens are isolated from other branches. This preserves independent rule evaluation within a single forward pass.

near-lossless acceleration for fast-mode multi-rule inference.

C.1 Motivation

The core challenge is to keep the decision format as close as possible to single-rule inference while removing its linear latency cost. Serial inference gives each rule an isolated context and therefore tends to be more reliable, but latency grows with the number of rules. Batch inference with frameworks such as vLLM can exploit GPU parallelism across independent sequences, but still redundantly encodes the shared image and prompt tokens for every rule. Packing all rules into one ordinary context removes this redundancy, but the model can attend across unrelated rule descriptions, which dilutes the evidence signal and causes missed detections. RI-Mask resolves these issues by sharing the content prefix across all rule branches while blocking cross-rule attention, so each branch behaves like an independent single-rule classifier.

C.2 Method

RI-Mask exploits a key structural property of rule-conditioned safety classification: all rules share the same *content prefix* (system prompt + image tokens + user text), differing only in the *rule suffix* and the corresponding *output tokens*. We construct a single packed sequence that concatenates the shared prefix with all N rule suffixes and their respective output slots, then apply a custom attention mask to enforce the following constraints:

1. **Shared prefix visibility:** All tokens in the content prefix (system instruction, image features, user text) are visible to every rule branch. This avoids redundant KV-cache computation for the dominant portion of the sequence.
2. **Rule isolation:** Each rule suffix and its output tokens can only attend to the shared prefix and to themselves. They *cannot* attend to other rule suffixes or their outputs, preventing cross-rule semantic interference.
3. **Causal generation:** Within each rule branch, output tokens follow standard causal (left-to-right) attention, ensuring autoregressive generation remains valid.

Formally, let the packed sequence be $S = [P; R_1; O_1; R_2; O_2; \dots; R_N; O_N]$, where P is the shared content prefix, R_i is the i -th rule description, and O_i is the output slot for rule i . The attention mask M is defined as:

$$M[j, k] = \begin{cases} 1 & \text{if } k \in P \text{ (prefix globally visible)} \\ 1 & \text{if } k \in (R_i \cup O_i) \text{ and } j \in (R_i \cup O_i) \text{ and } k \leq j \\ 0 & \text{otherwise} \end{cases} \quad (10)$$

This mask structure is analogous to the *message-tree encoding* scheme proposed in Molmo2 (Deitke et al., 2026), where a shared visual prefix is attended to by multiple independent annotation branches. In our setting, the “tree” has a single root (the content prefix) with N leaf branches (one per rule), each producing an independent safety judgment.

C.3 Implementation

We implement RI-Mask using PyTorch’s FlexAttention API, which supports arbitrary block-sparse attention patterns without requiring custom CUDA kernels. The position encoding is modified so that each rule branch

continues the positional IDs from the end of the shared prefix, ensuring that all branches share consistent positional semantics relative to the content. At inference time, the model generates all N rule outputs in parallel within a single forward pass, and the results are split back into per-rule predictions.

C.4 Discussion

The RI-Mask approach is particularly effective for SingGuard’s deployment scenario because most online requests are handled by fast-mode rule checking rather than slow deliberation. In this regime, the content prefix (image tokens + system prompt) dominates the total sequence length, rule descriptions are relatively short, and each branch only needs to emit a compact hit / not-hit output. RI-Mask therefore avoids repeatedly encoding the same content while preserving rule isolation, giving near-lossless parallel acceleration compared with independent single-rule inference. For slow-mode reasoning where outputs are longer, RI-Mask can still be applied but with reduced parallelism to fit within GPU memory constraints.

This technique generalizes to any multi-label or multi-aspect evaluation scenario where a shared input must be independently assessed against multiple criteria—a common pattern in content moderation, compliance checking, and multi-dimensional quality assessment.